CENG 506 Deep Learning

Lecture 3 – Neural Networks and Backpropagation

Slides were prepared using the course material of Stanford's Machine Learning Course (CS229 by Andrew Ng) and CNN Course (CS231n by Fei-Fei, Karpathy, Johnson)

Brain neurons

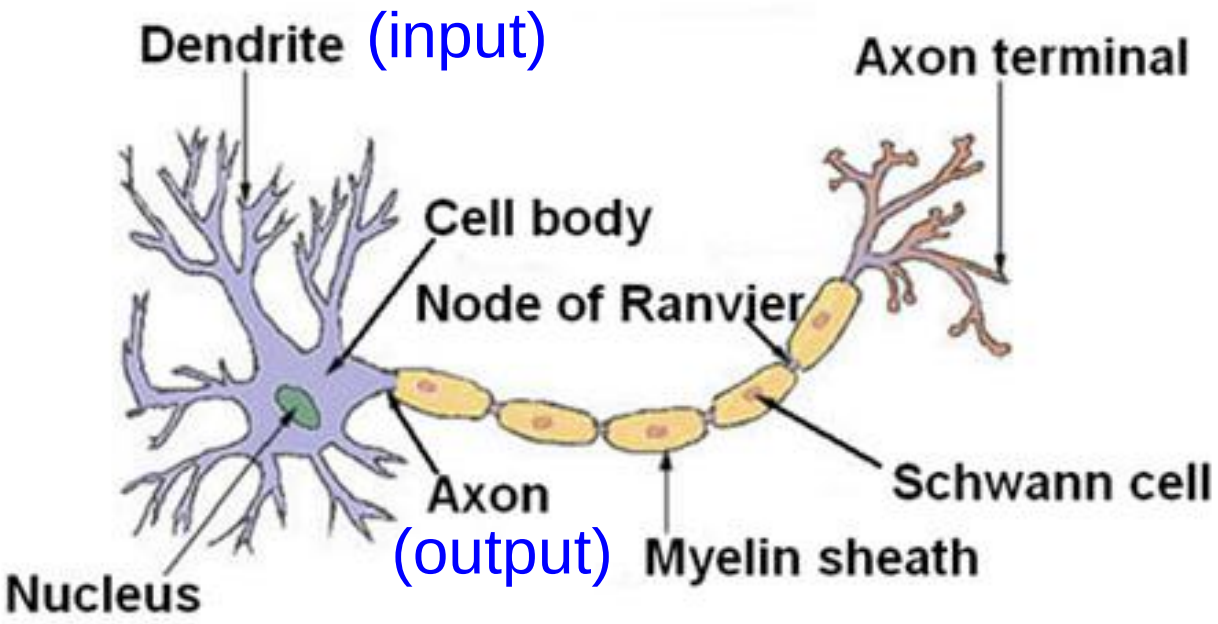

Neuron Model

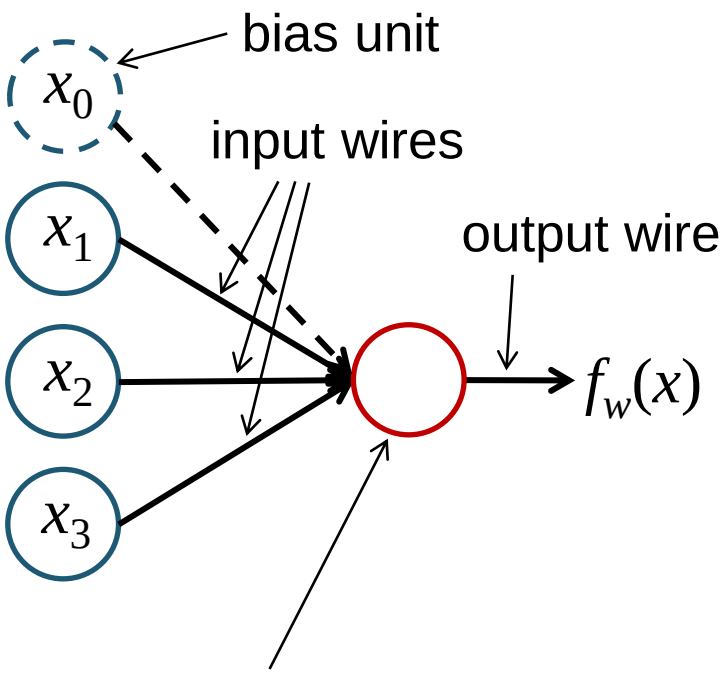

$$f_w(x) = \frac{1}{1+e^{-w^T x}}$$

$$x = \begin{bmatrix} x_0 \\ x_1 \\ x_2 \\ x_3 \end{bmatrix} \qquad w = \begin{bmatrix} w_0 \\ w_1 \\ w_2 \\ w_3 \end{bmatrix}$$

The parameters *w* are called ‘weights’.

This computation is determined by the *activation function*.
In this example, activation function is sigmoid function.

Example: AND

Let's build a linear classifier with a one-neuron model (no hidden layers).

Logical AND operation

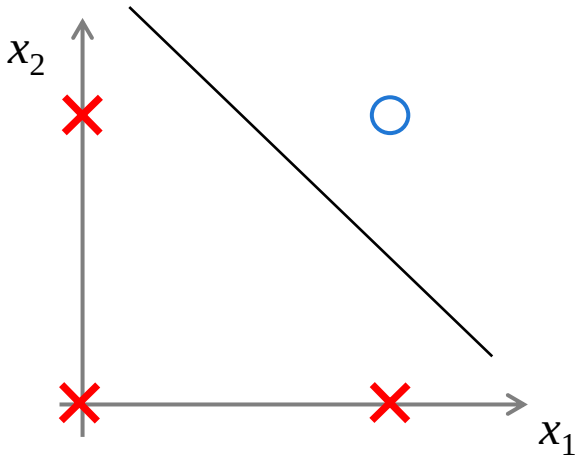

x_1, x_2 are binary (0 or 1).

y=1 if x_1 AND x_2

(So blue circle is the positive class)

Example: AND

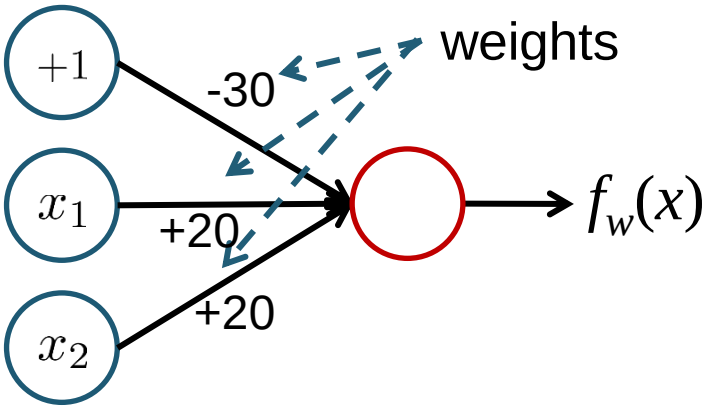

g: sigmoid function

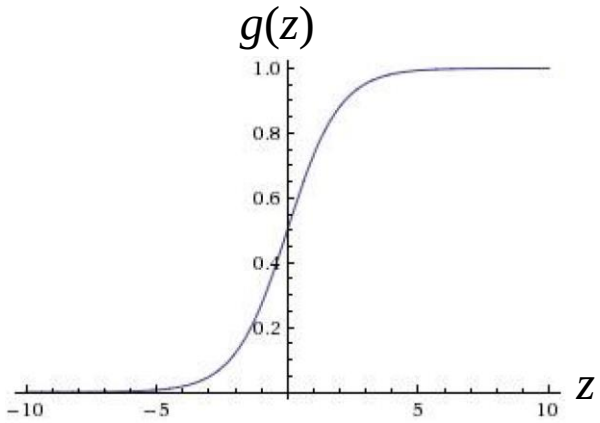

$$f_w(x) = g(-30+20x_1+20x_2)$$

x_1	x_2	$f_w(x)$
0	0	$g(-30) \approx 0$
0	1	$g(-10) \approx 0$
1	0	$g(-10) \approx 0$
1	1	$g(10) \approx 1$

$$f_w(x) \approx x_1 \text{ AND } x_2$$

Example: OR

Another linear classifier for logical OR operation

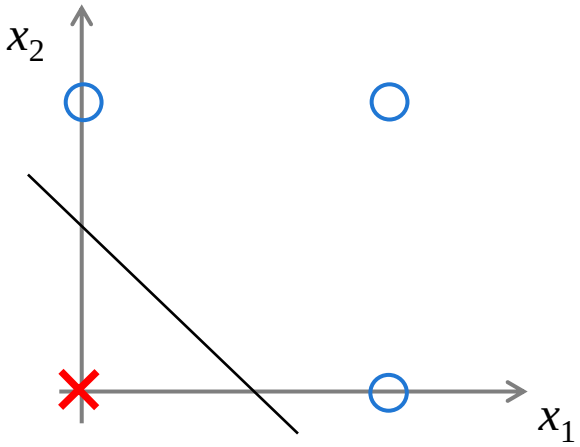

x_1, x_2 are binary (0 or 1).
y=1 if x_1 OR x_2

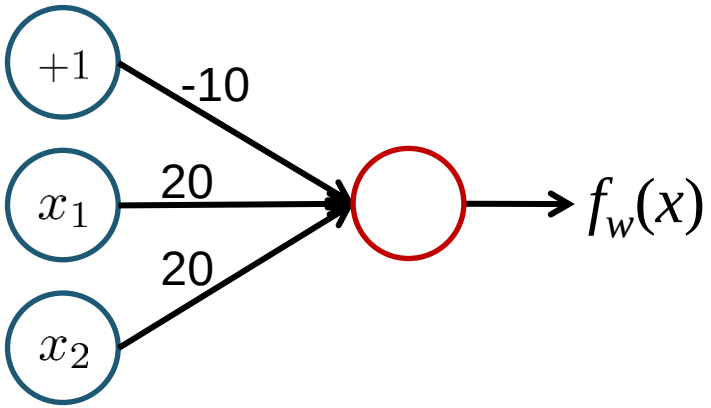

x_1	x_2	$f_w(x)$
0	0	$g(-10) \approx 0$
0	1	$g(10) \approx 1$
1	0	$g(10) \approx 1$
1	1	$g(30) \approx 1$

$f_w(x) = g(-10+20x_1+20x_2)$

$f_w(x) \approx x_1$ OR x_2

Nonlinear classification example

Wait a minute! NNs are good for nonlinear classification.

This is done by adding more layers.

A non-linear classification example: XNOR

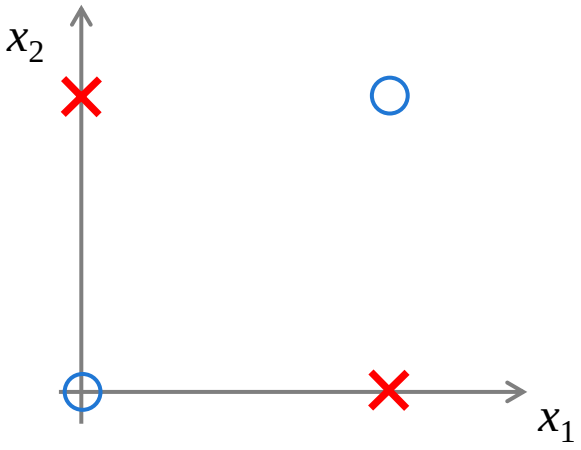

y=0 if x_1 XOR x_2

y=1 if NOT (x_1 XOR x_2)= x_1 XNOR x_2

(blue circles are the positive class)

Hint: x_1 XNOR x_2 = (NOT x_1 AND NOT x_2) OR (x_1 AND x_2)

Example: XNOR

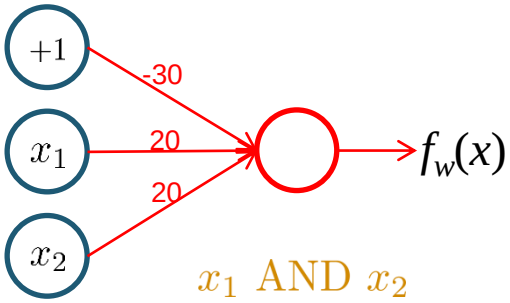

x_1 AND x_2

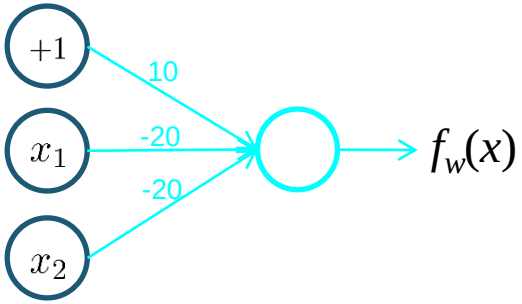

(NOT x_1) AND (NOT x_2)

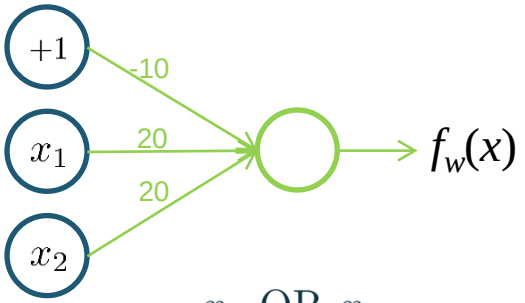

x_1 OR x_2

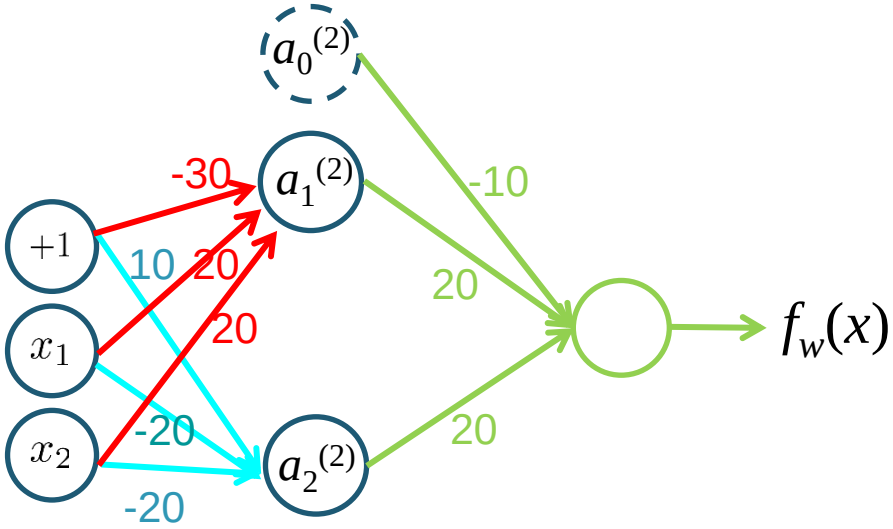

x_1	x_2	$a_1^{(2)}$	$a_2^{(2)}$	$f_w(x)$
0	0	0	1	1
0	1	0	0	0
1	0	0	0	0
1	1	1	0	1

2-layer Neural Network (or 1-hidden-layer neural network)

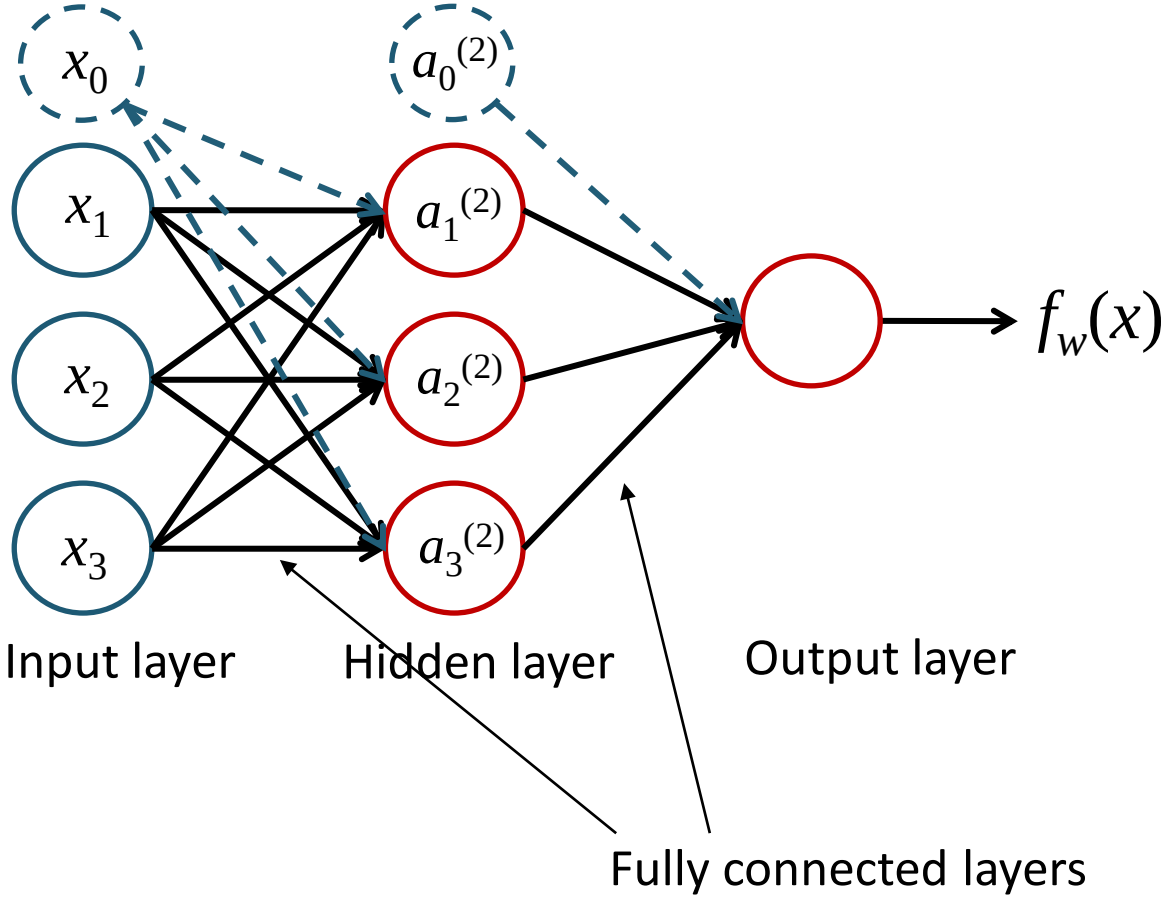

3-layer Neural Network (or 2-hidden-layer neural network)

As we go further in layers, more complex functions are modeled.

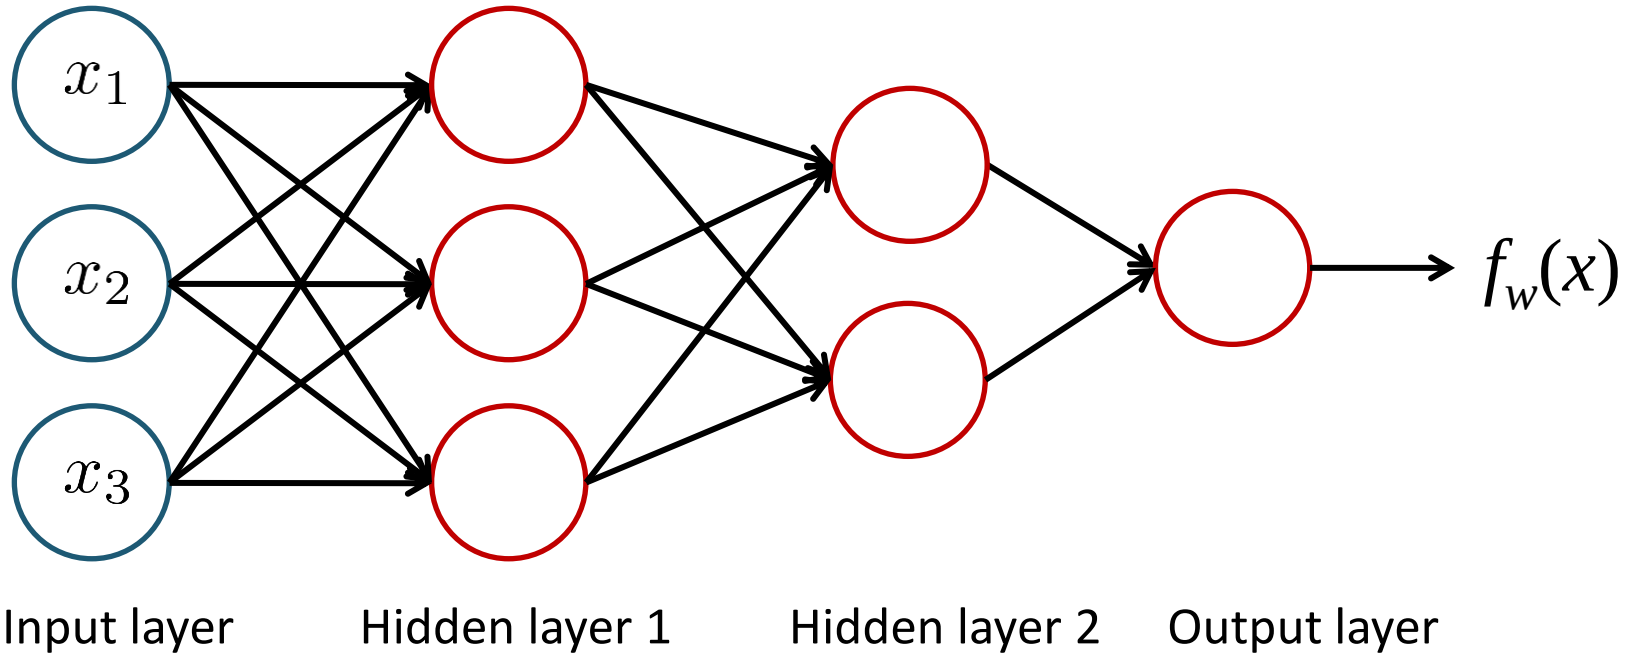

Effect of the number of hidden neurons/layers

More neurons, more capacity to learn complex boundaries

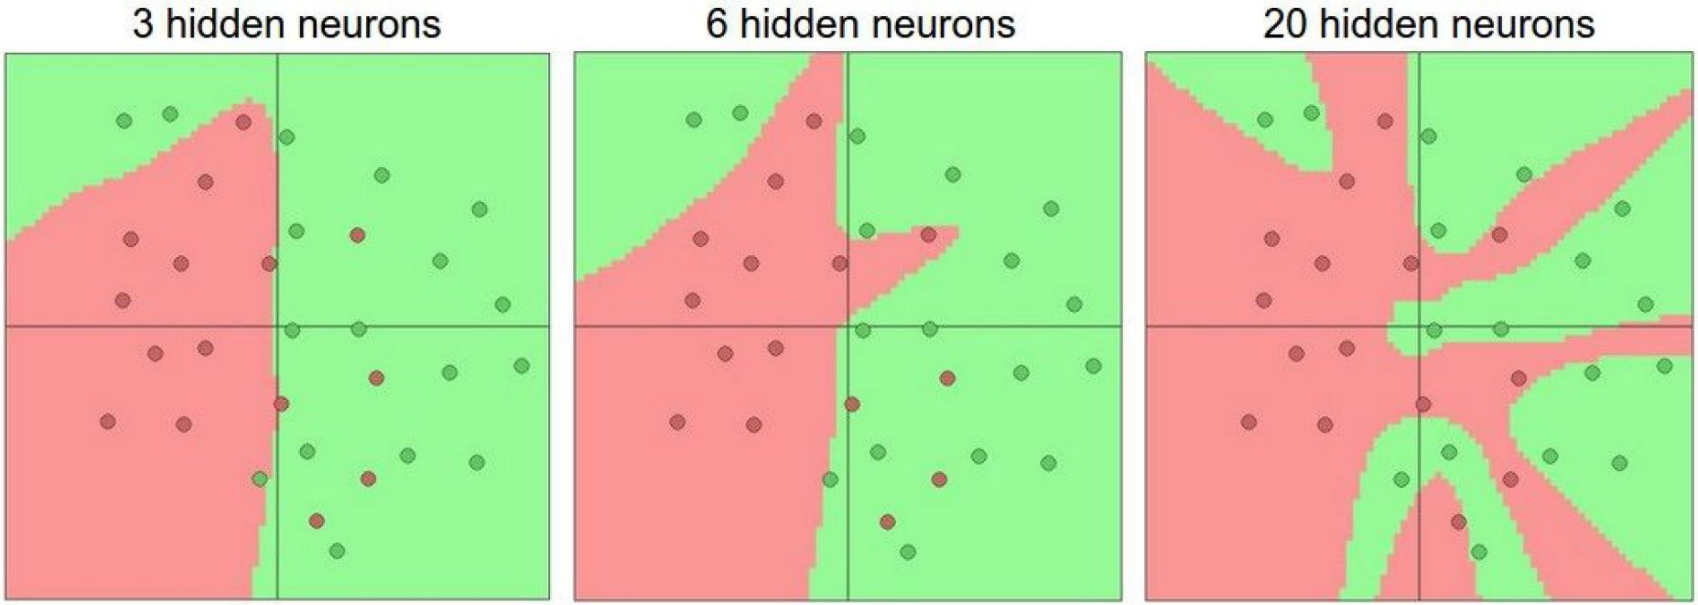

Forward Propagation

$$z_1^{(2)} = w_{10}^{(1)} + w_{11}^{(1)} x_1 + w_{12}^{(1)} x_2 + w_{13}^{(1)} x_3$$

$$a_1^{(2)} = g(z_1^{(2)})$$

Remember g is the sigmoid function

$a_i^{(j)}$ = "activation" of unit i in layer j

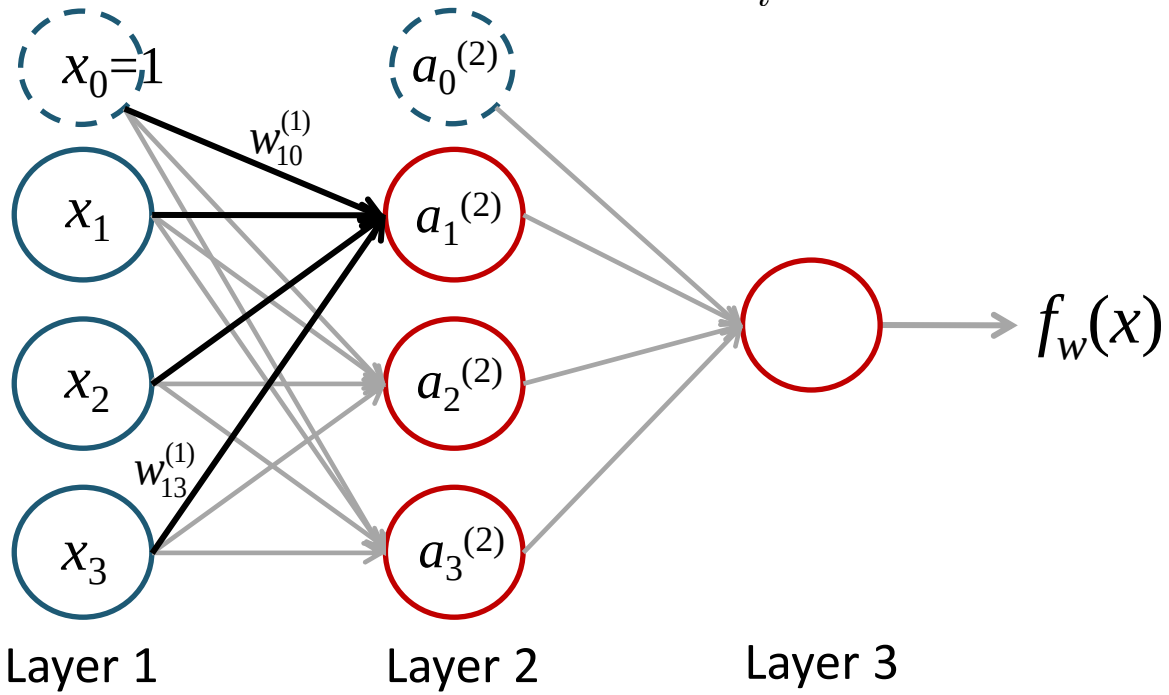

Forward Propagation

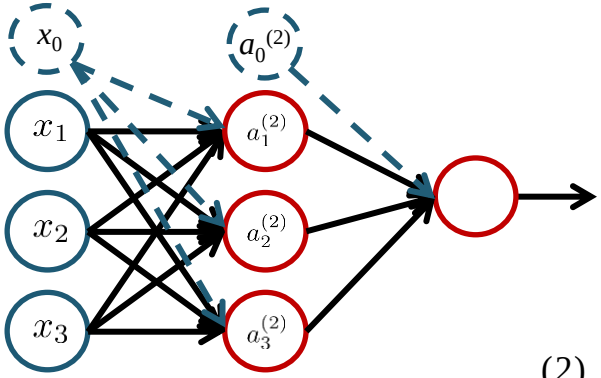

$$a_1^{(2)} = g(\underbrace{w_{10}^{(1)}x_0 + w_{11}^{(1)}x_1 + w_{12}^{(1)}x_2 + w_{13}^{(1)}x_3}_{z_1^{(2)}})$$

$$a_2^{(2)} = g(\underbrace{w_{20}^{(1)}x_0 + w_{21}^{(1)}x_1 + w_{22}^{(1)}x_2 + w_{23}^{(1)}x_3}_{z_2^{(2)}})$$

$$a_3^{(2)} = g(\underbrace{w_{30}^{(1)}x_0 + w_{31}^{(1)}x_1 + w_{32}^{(1)}x_2 + w_{33}^{(1)}x_3}_{z_3^{(2)}})$$

$$f_w(x) = a_1^{(3)} = g(\underbrace{w_{10}^{(2)}a_0^{(2)} + w_{11}^{(2)}a_1^{(2)} + w_{12}^{(2)}a_2^{(2)} + w_{13}^{(2)}a_3^{(2)}}_{z_1^{(3)}})$$

Forward Propagation

Important note:

If we do not use a non-linear activation function,

$$a_1^{(2)} = w_{10}^{(1)} + w_{11}^{(1)}x_1 + w_{12}^{(1)}x_2 + w_{13}^{(1)}x_3 \qquad a^{(2)} = W^{(1)}x \qquad f_w(x) = W^{(2)}W^{(1)}x$$

it doesn't work. Everything boils down to a linear product.

That's why we add non-linearity: $a_1^{(2)} = g(z_1^{(2)})$

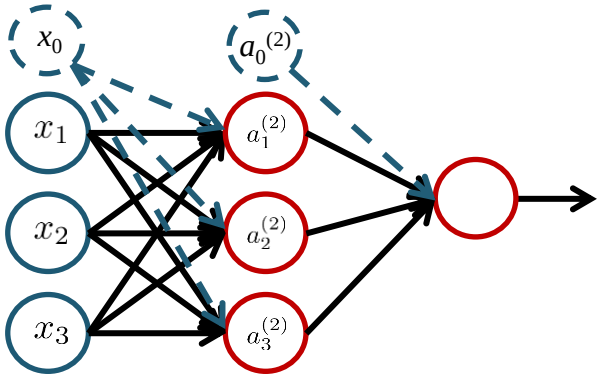

Activation functions

Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$

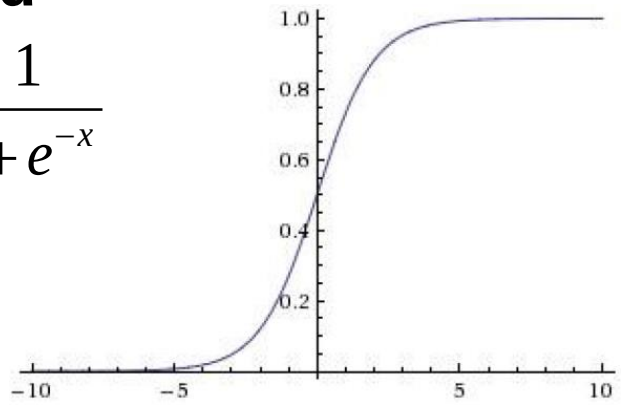

ReLU

$$\max(0, x)$$

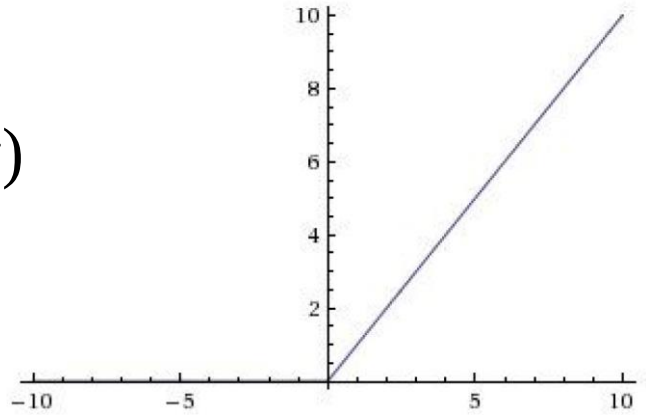

tanh

$$\tanh(x)$$

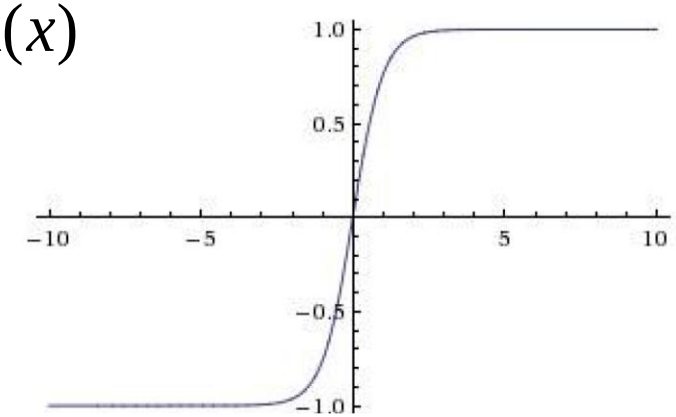

Leaky ReLU

$$\max(0.1x, x)$$

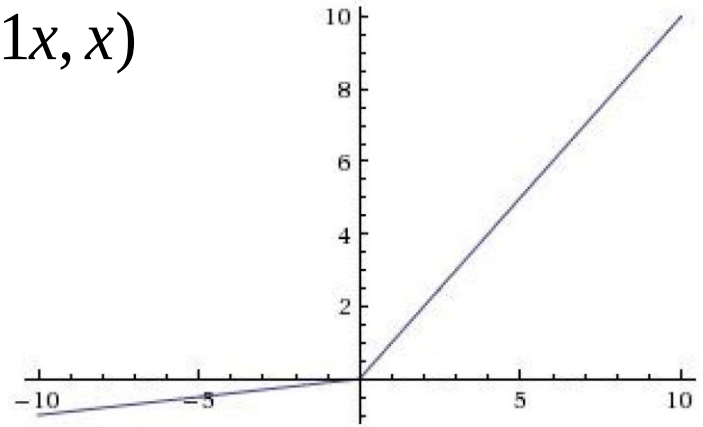

Neural Network Learning

We have learned about:

- Analogy with the human brain
- Layers of neural networks
- Forward propagation

But how do neural networks ‘learn’ actually?

Learning corresponds to determining the ‘weights’ of units.

These weights are optimized using a cost function.

We define ‘loss’ here.

Loss

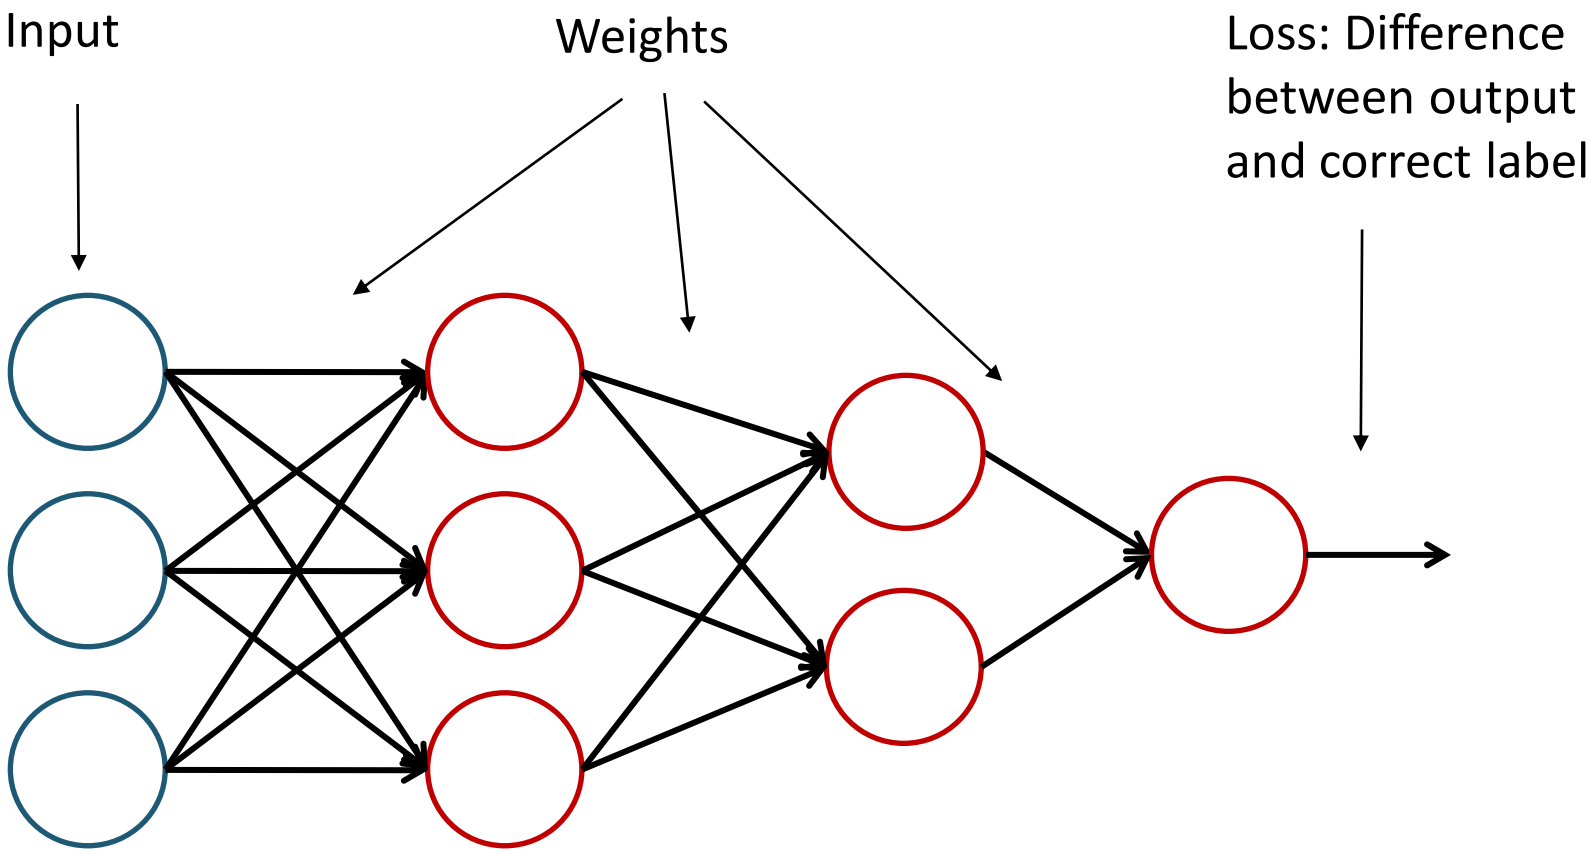

To update the weights, we ‘backpropagate’ the loss to previous layers.

Derivatives for backpropagation

$f(x, y, z) = (x + y)z$

E.g. x=-2, y=5, z=-4

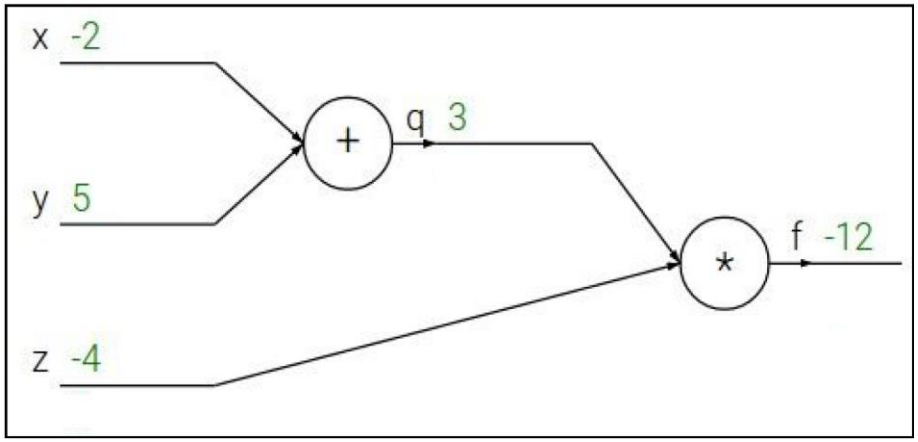

$$q = x + y \qquad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \qquad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want: $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$

Derivatives for backpropagation

$$f(x, y, z) = (x + y)z$$

E.g. x=-2, y=5, z=-4

$$q = x + y \qquad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \qquad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

Want: $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$

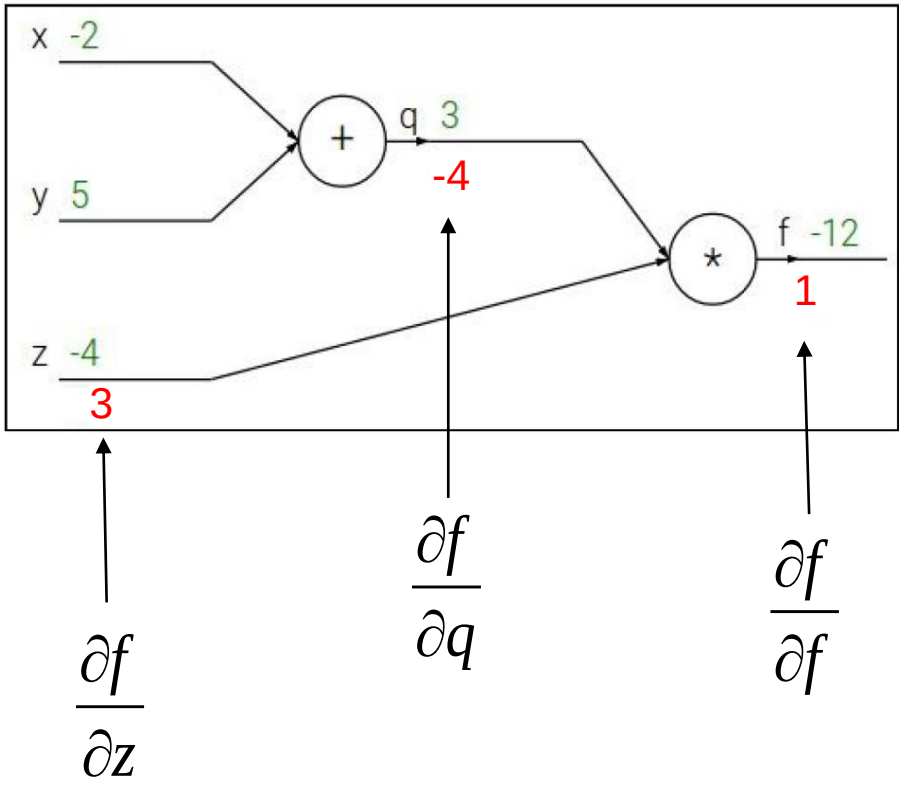

Derivatives for backpropagation

$f(x, y, z) = (x + y)z$

E.g. x=-2, y=5, z=-4

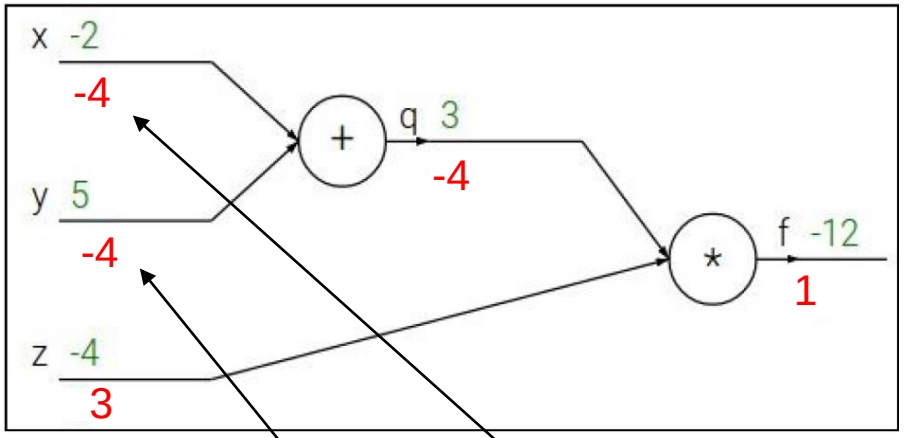

$$q = x + y \qquad \frac{\partial q}{\partial x} = 1, \frac{\partial q}{\partial y} = 1$$

$$f = qz \qquad \frac{\partial f}{\partial q} = z, \frac{\partial f}{\partial z} = q$$

$\frac{\partial f}{\partial y}$ $\frac{\partial f}{\partial x}$

Want: $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$

Chain rule: $\frac{\partial f}{\partial y} = \frac{\partial f}{\partial q}\frac{\partial q}{\partial y}$ $\frac{\partial f}{\partial x} = \frac{\partial f}{\partial q}\frac{\partial q}{\partial x}$

Gradients for backpropagation

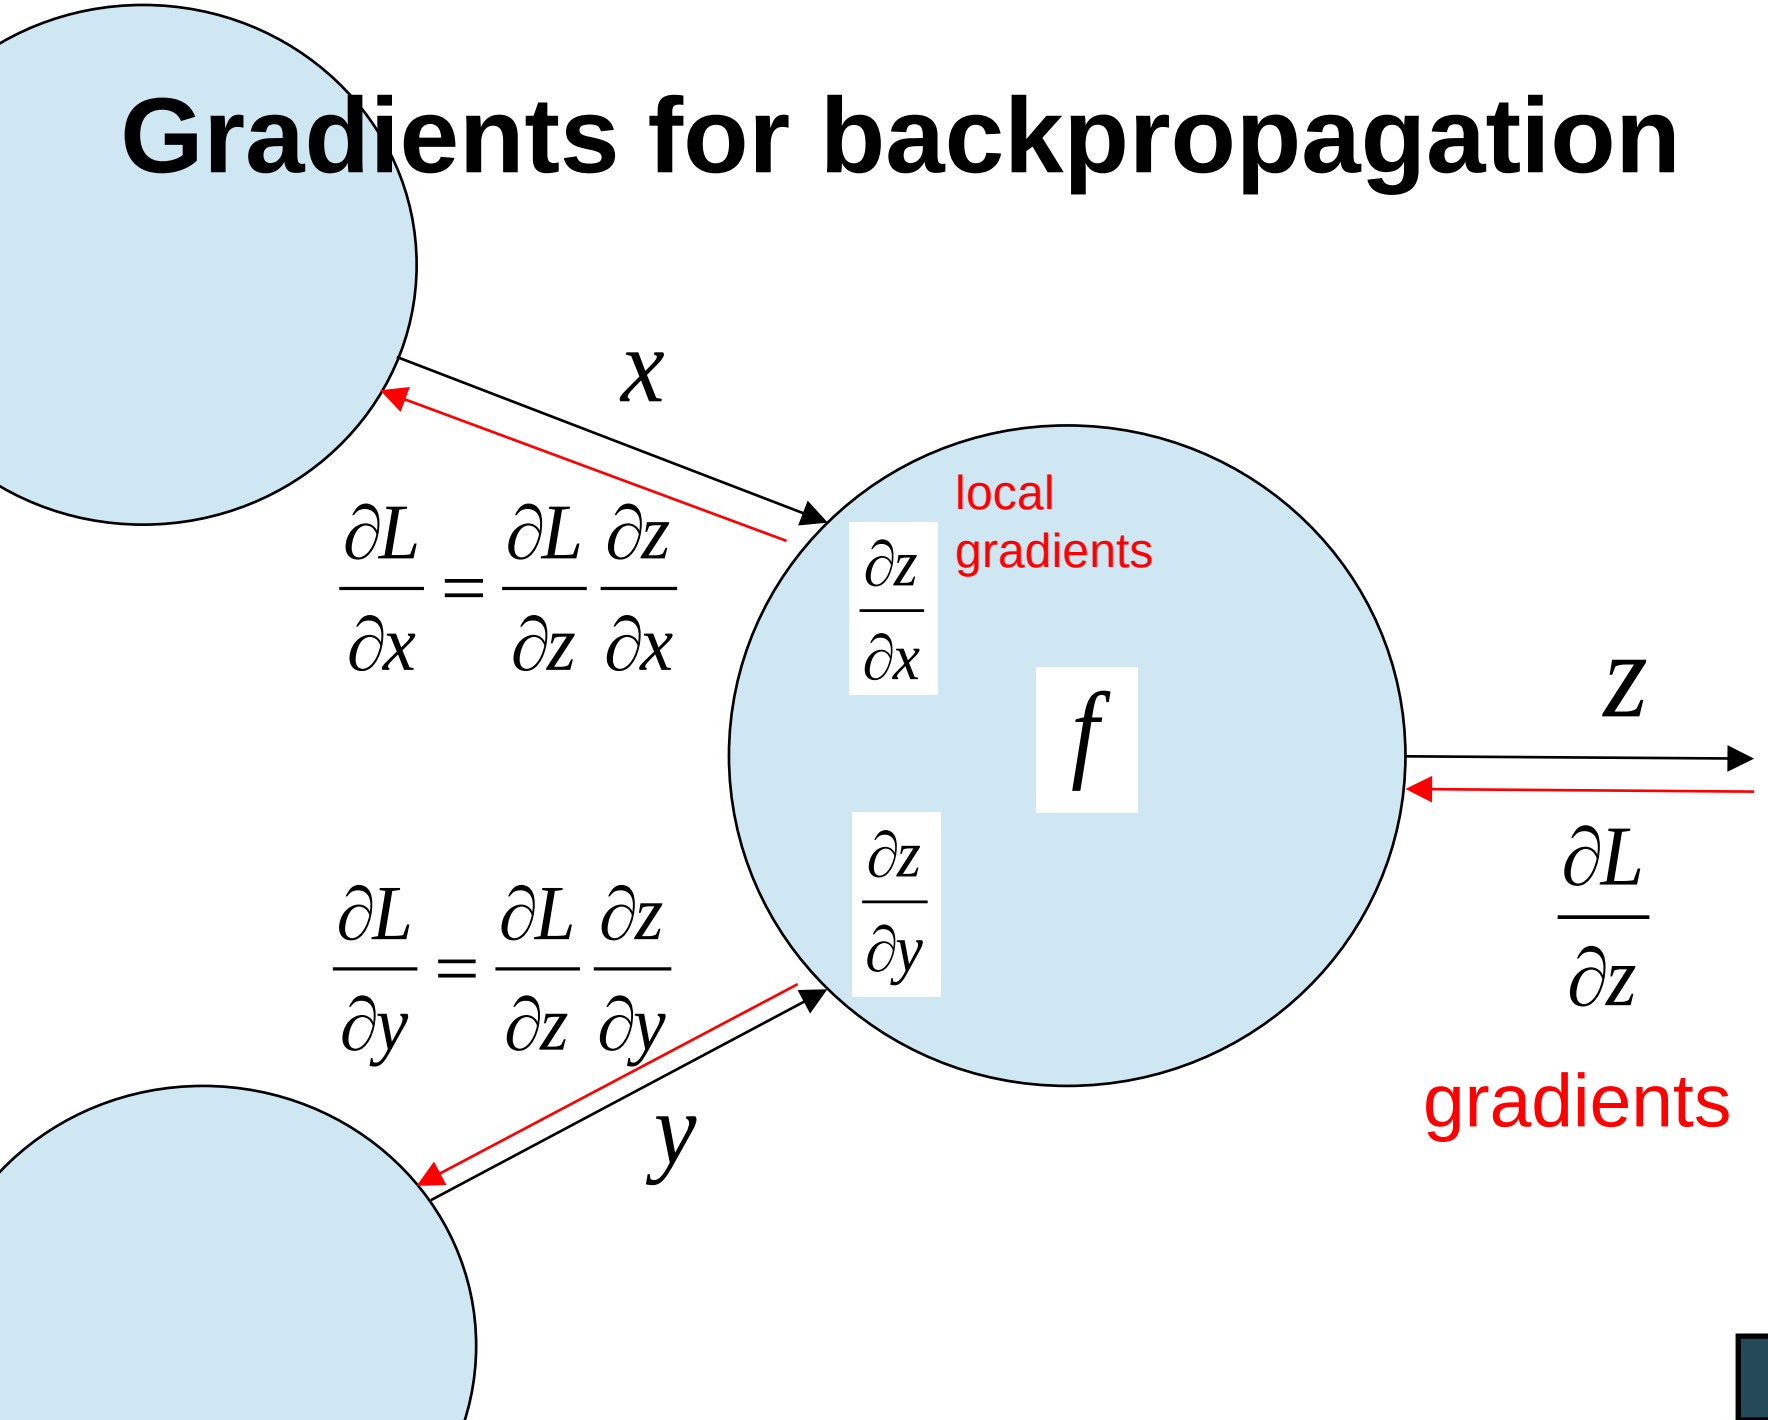

$$\frac{\partial L}{\partial x} = \frac{\partial L}{\partial z}\frac{\partial z}{\partial x}$$

$$\frac{\partial L}{\partial y} = \frac{\partial L}{\partial z}\frac{\partial z}{\partial y}$$

Another example

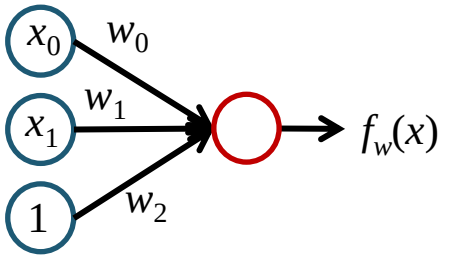

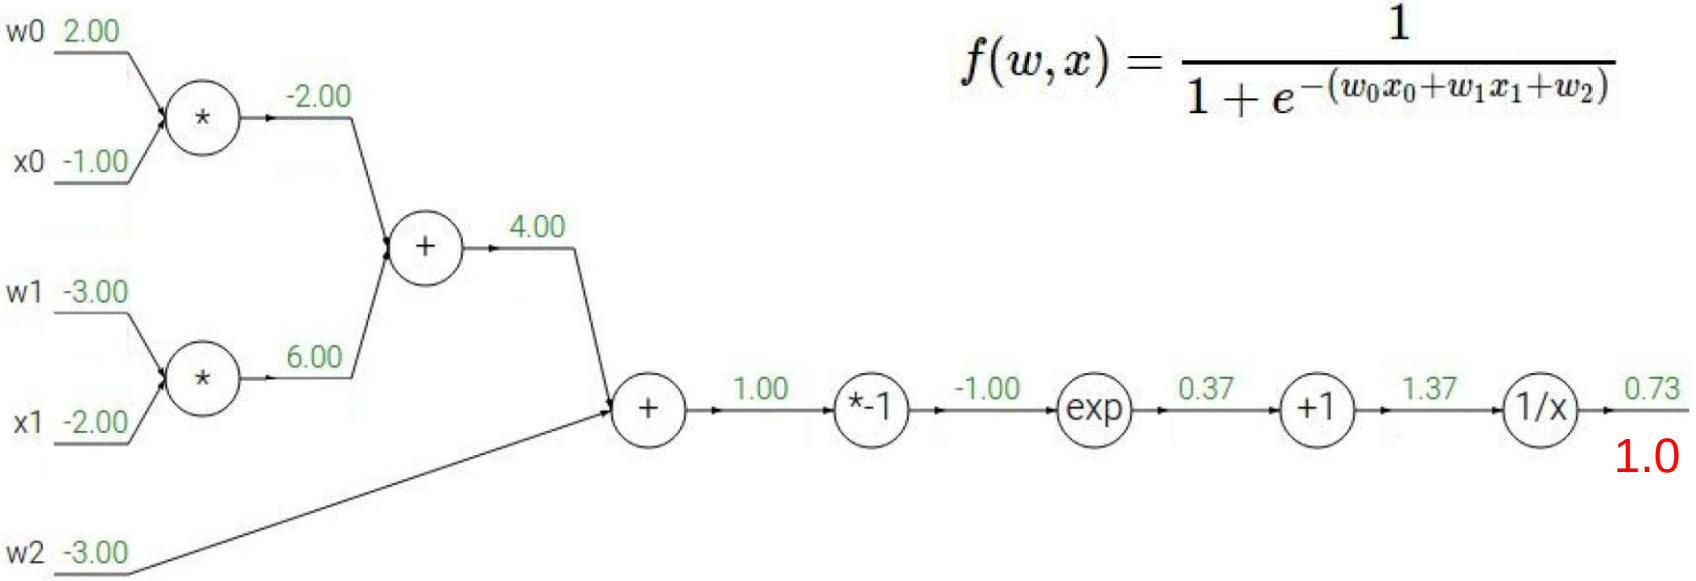

$$f(w,x) = \frac{1}{1+e^{-(w_0x_0+w_1x_1+w_2)}}$$

$$f(x) = e^x \quad \rightarrow \quad \frac{df}{dx} = e^x$$

$$f_a(x) = ax \quad \rightarrow \quad \frac{df}{dx} = a$$

$$f(x) = \frac{1}{x} \quad \rightarrow \quad \frac{df}{dx} = -1/x^2$$

$$f_c(x) = c + x \quad \rightarrow \quad \frac{df}{dx} = 1$$

Another example

$$f(w,x) = \frac{1}{1+e^{-(w_0x_0+w_1x_1+w_2)}}$$

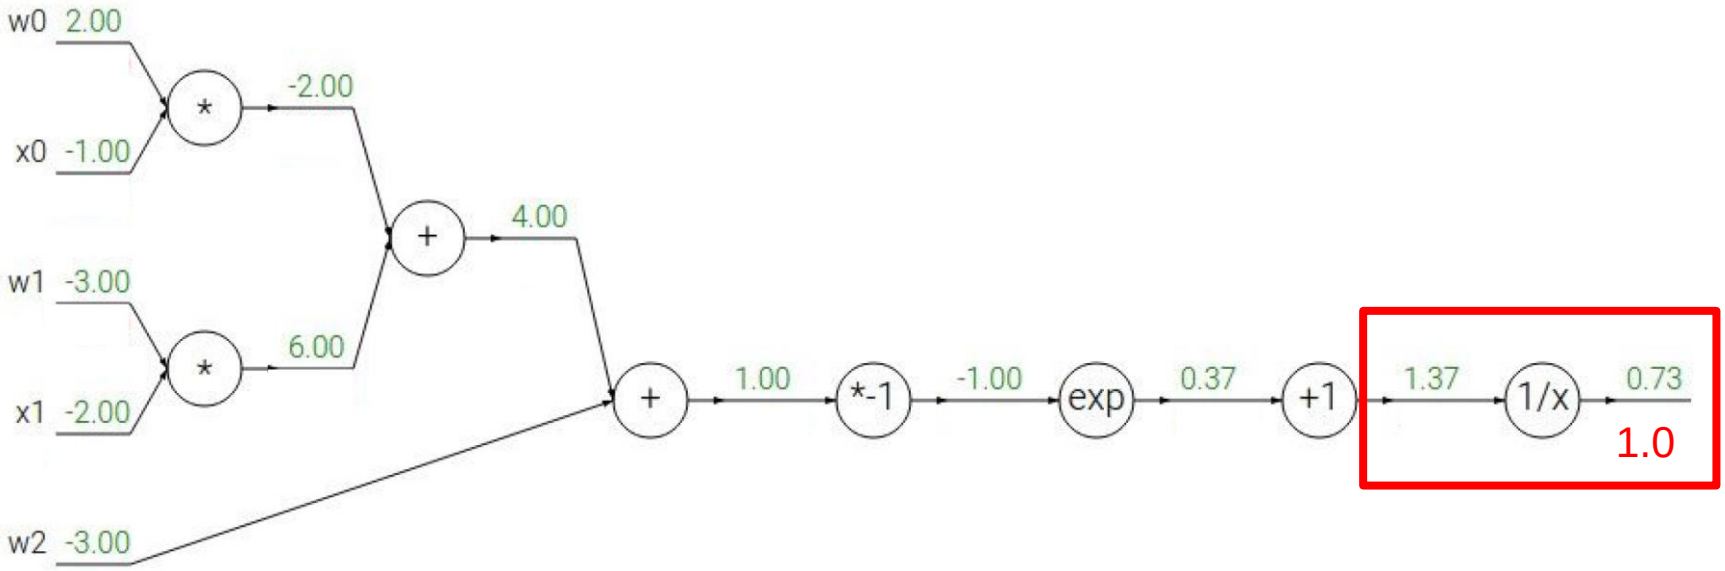

$$f(x) = e^x \quad \rightarrow \quad \frac{df}{dx} = e^x$$

$$f_a(x) = ax \quad \rightarrow \quad \frac{df}{dx} = a$$

$$f(x) = \frac{1}{x} \quad \rightarrow \quad \frac{df}{dx} = -1/x^2$$

$$f_c(x) = c + x \quad \rightarrow \quad \frac{df}{dx} = 1$$

Another example

$$f(w,x) = \frac{1}{1+e^{-(w_0x_0+w_1x_1+w_2)}}$$

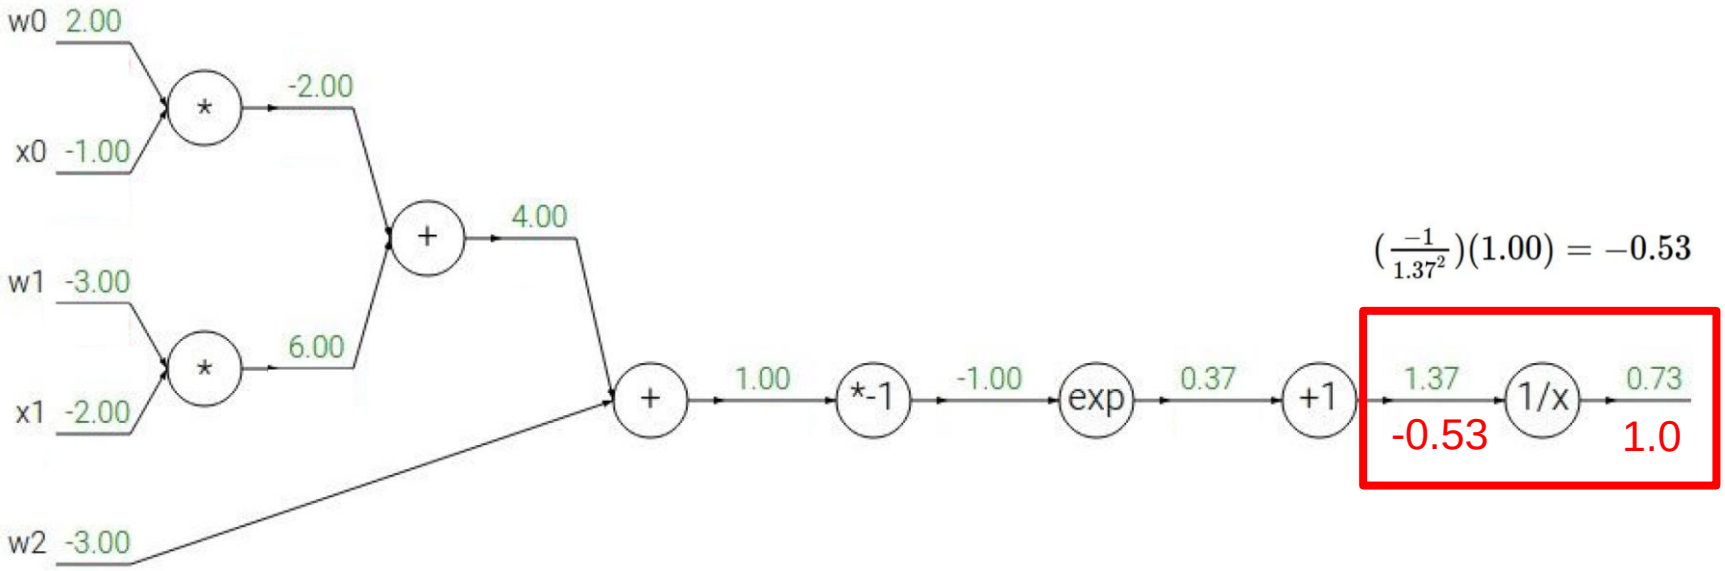

$f(x) = e^x \rightarrow \frac{df}{dx} = e^x$

$f_a(x) = ax \rightarrow \frac{df}{dx} = a$

$f(x) = \frac{1}{x} \rightarrow \frac{df}{dx} = -1/x^2$

$f_c(x) = c + x \rightarrow \frac{df}{dx} = 1$

Another example

$$f(w,x) = \frac{1}{1+e^{-(w_0x_0+w_1x_1+w_2)}}$$

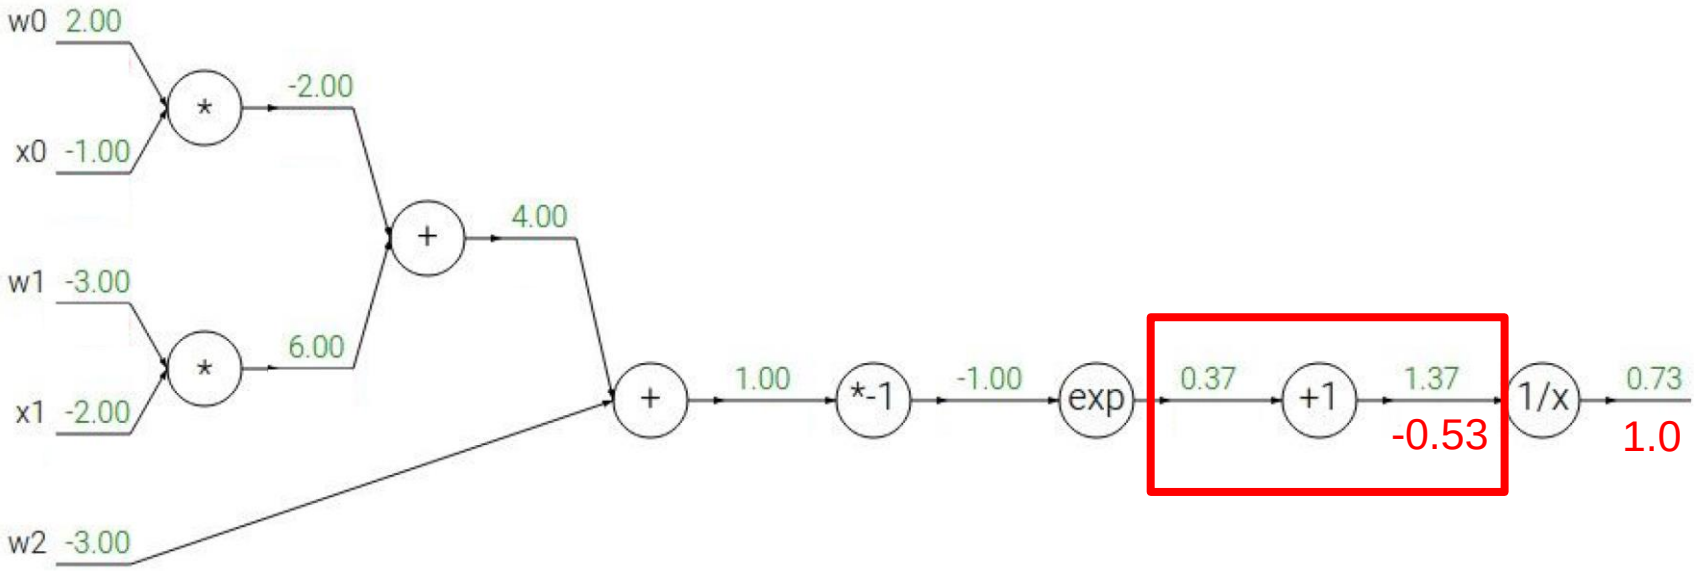

$$f(x) = e^x \rightarrow \frac{df}{dx} = e^x$$

$$f_a(x) = ax \rightarrow \frac{df}{dx} = a$$

$$f(x) = \frac{1}{x} \rightarrow \frac{df}{dx} = -1/x^2$$

$$f_c(x) = c + x \rightarrow \frac{df}{dx} = 1$$

Another example

$$f(w,x) = \frac{1}{1+e^{-(w_0x_0+w_1x_1+w_2)}}$$

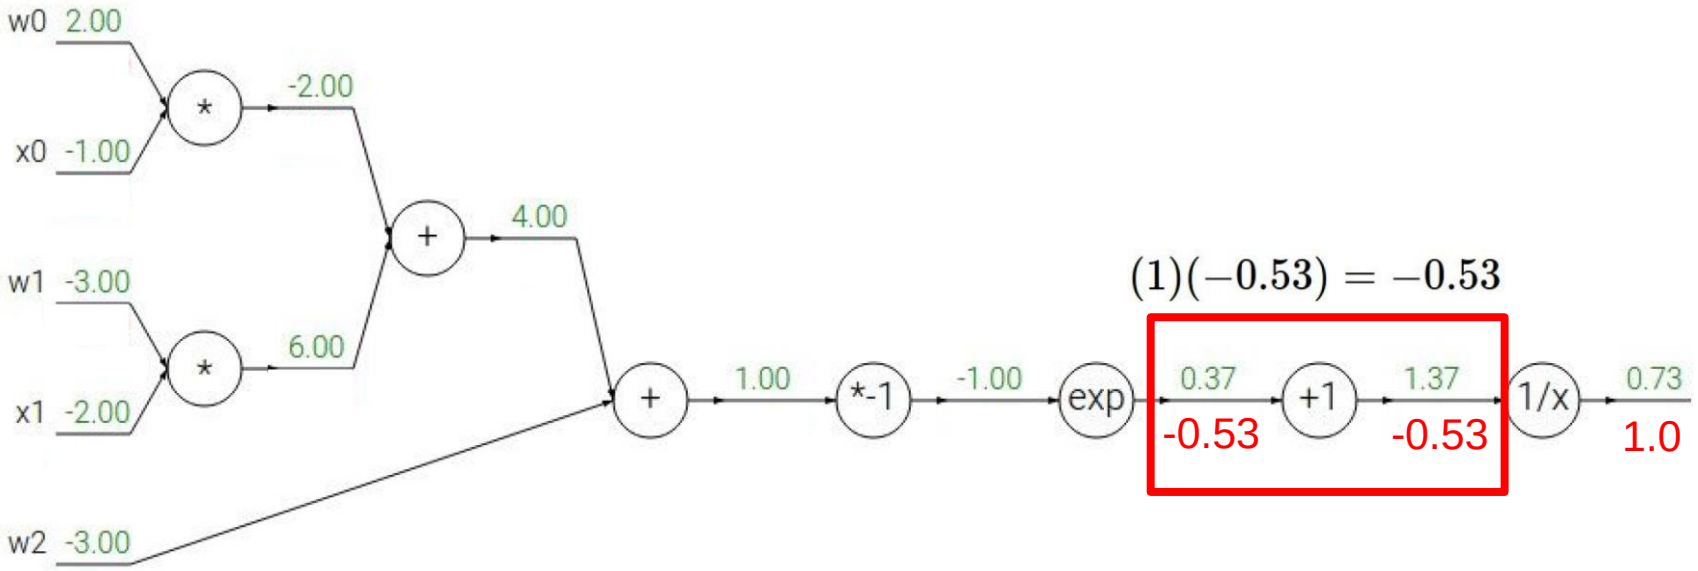

$$f(x) = e^x \rightarrow \frac{df}{dx} = e^x$$

$$f_a(x) = ax \rightarrow \frac{df}{dx} = a$$

$$f(x) = \frac{1}{x} \rightarrow \frac{df}{dx} = -1/x^2$$

$$f_c(x) = c + x \rightarrow \frac{df}{dx} = 1$$

Another example

$$f(w,x) = \frac{1}{1+e^{-(w_0x_0+w_1x_1+w_2)}}$$

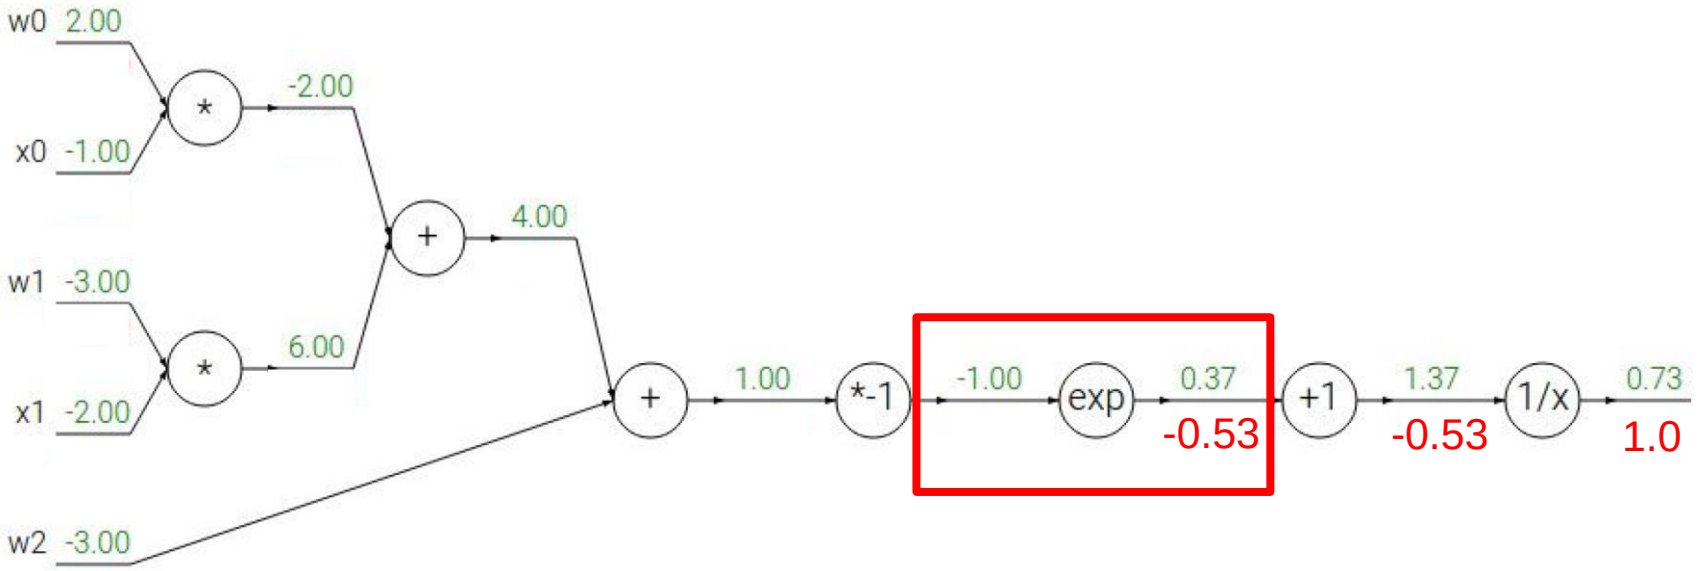

$$f(x) = e^x \rightarrow \frac{df}{dx} = e^x$$

$$f_a(x) = ax \rightarrow \frac{df}{dx} = a$$

$$f(x) = \frac{1}{x} \rightarrow \frac{df}{dx} = -1/x^2$$

$$f_c(x) = c + x \rightarrow \frac{df}{dx} = 1$$

Another example

$$f(w,x) = \frac{1}{1+e^{-(w_0x_0+w_1x_1+w_2)}}$$

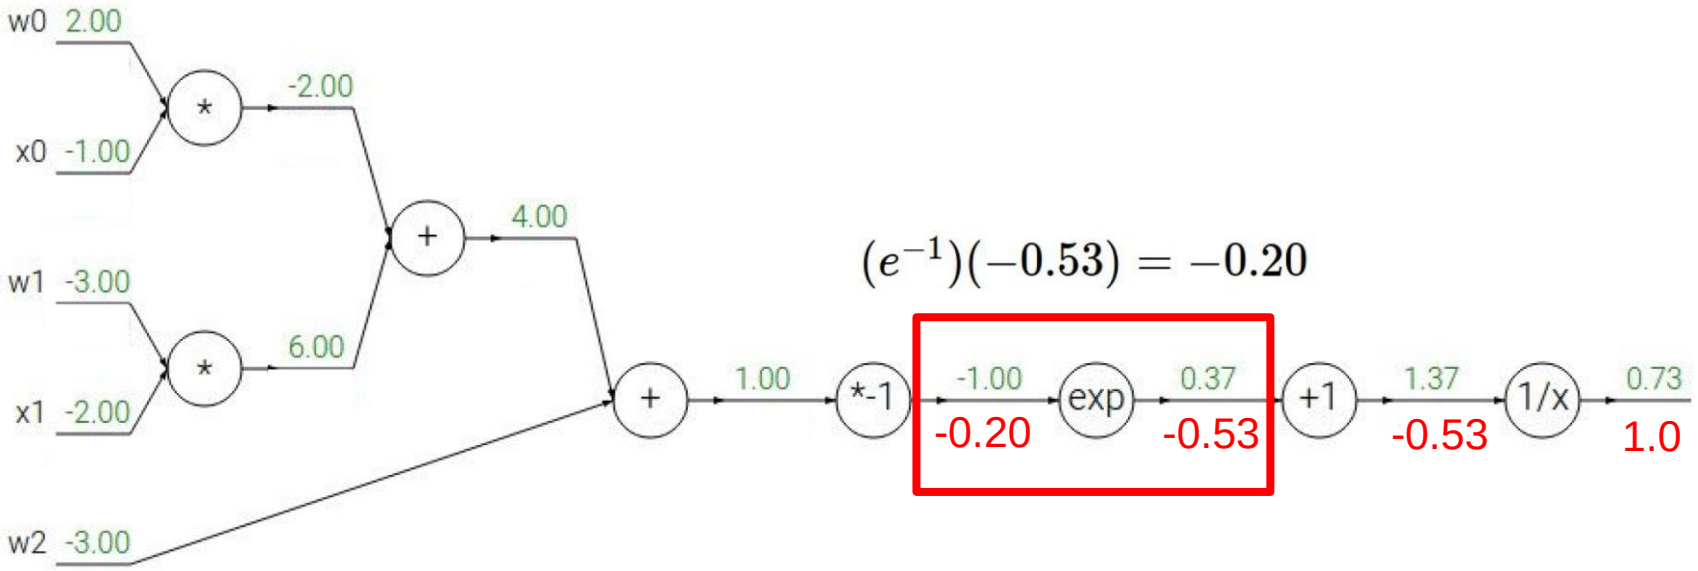

$$(e^{-1})(-0.53) = -0.20$$

$$f(x) = e^x \rightarrow \frac{df}{dx} = e^x$$

$$f_a(x) = ax \rightarrow \frac{df}{dx} = a$$

$$f(x) = \frac{1}{x} \rightarrow \frac{df}{dx} = -1/x^2$$

$$f_c(x) = c + x \rightarrow \frac{df}{dx} = 1$$

Another example

$$f(w,x) = \frac{1}{1+e^{-(w_0x_0+w_1x_1+w_2)}}$$

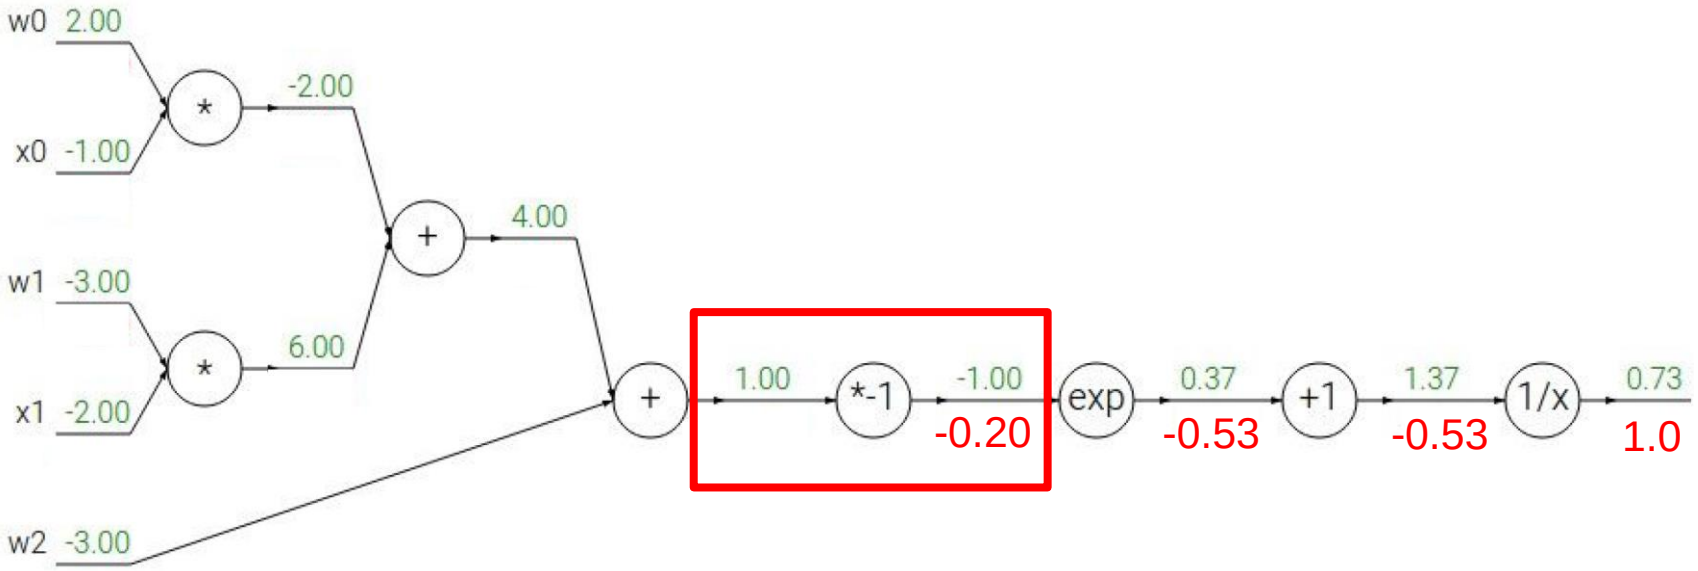

$$f(x) = e^x \rightarrow \frac{df}{dx} = e^x$$

$$f_a(x) = ax \rightarrow \frac{df}{dx} = a$$

$$f(x) = \frac{1}{x} \rightarrow \frac{df}{dx} = -1/x^2$$

$$f_c(x) = c + x \rightarrow \frac{df}{dx} = 1$$

Another example

$$f(w,x) = \frac{1}{1+e^{-(w_0x_0+w_1x_1+w_2)}}$$

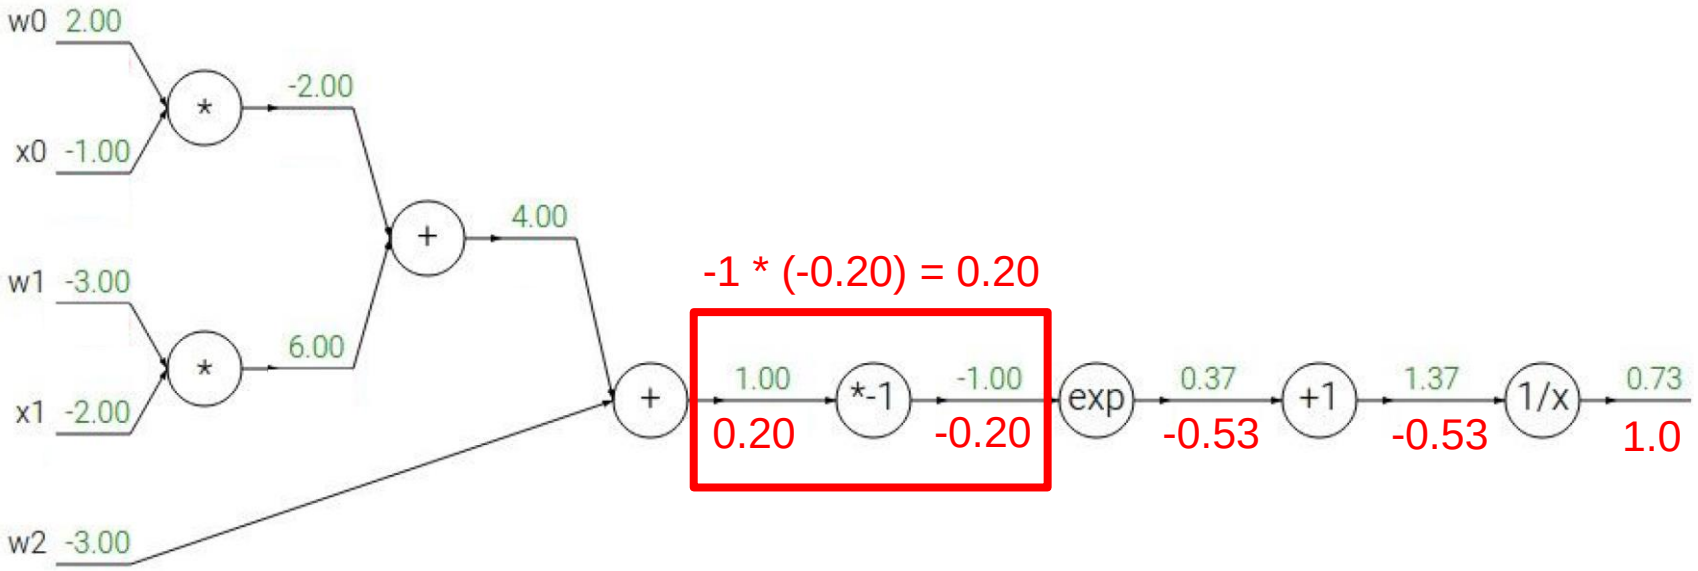

$$f(x) = e^x \rightarrow \frac{df}{dx} = e^x$$

$$f_a(x) = ax \rightarrow \frac{df}{dx} = a$$

$$f(x) = \frac{1}{x} \rightarrow \frac{df}{dx} = -1/x^2$$

$$f_c(x) = c + x \rightarrow \frac{df}{dx} = 1$$

Another example

$$f(w,x) = \frac{1}{1+e^{-(w_0x_0+w_1x_1+w_2)}}$$

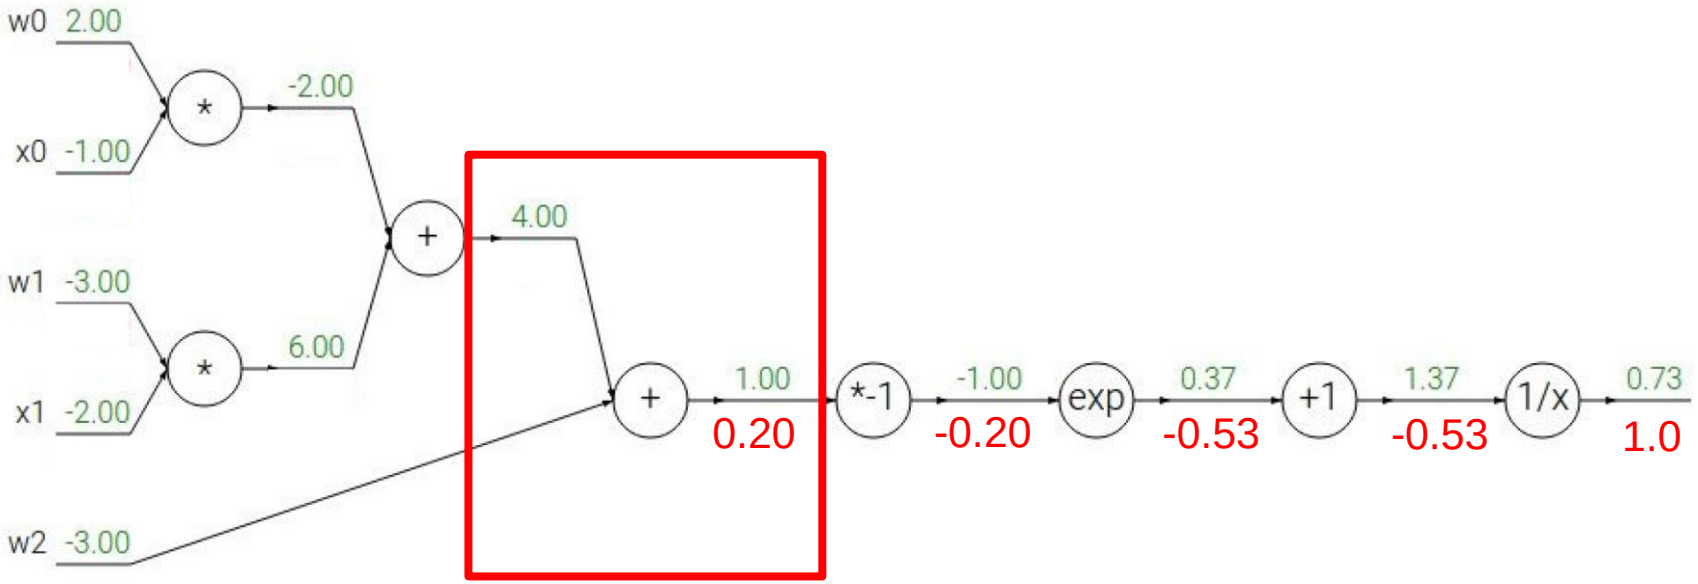

$$f(x) = e^x \quad \rightarrow \quad \frac{df}{dx} = e^x$$

$$f_a(x) = ax \quad \rightarrow \quad \frac{df}{dx} = a$$

$$f(x) = \frac{1}{x} \quad \rightarrow \quad \frac{df}{dx} = -1/x^2$$

$$f_c(x) = c + x \quad \rightarrow \quad \frac{df}{dx} = 1$$

Another example

$$f(w,x) = \frac{1}{1+e^{-(w_0x_0+w_1x_1+w_2)}}$$

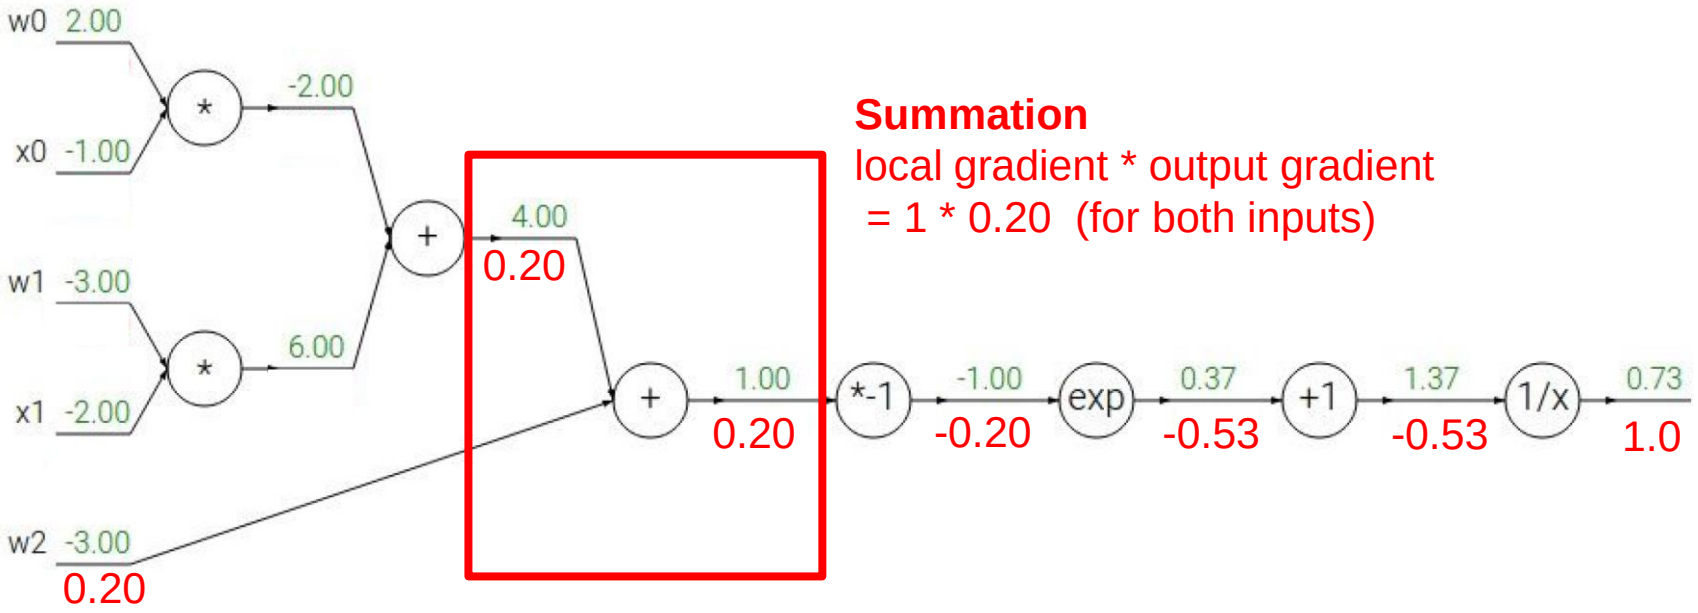

Summation
local gradient * output gradient
= 1 * 0.20 (for both inputs)

$$f(x) = e^x \rightarrow \frac{df}{dx} = e^x \qquad f(x) = \frac{1}{x} \rightarrow \frac{df}{dx} = -1/x^2$$

$$f_a(x) = ax \rightarrow \frac{df}{dx} = a \qquad f_c(x) = c + x \rightarrow \frac{df}{dx} = 1$$

Another example

$$f(w,x) = \frac{1}{1+e^{-(w_0x_0+w_1x_1+w_2)}}$$

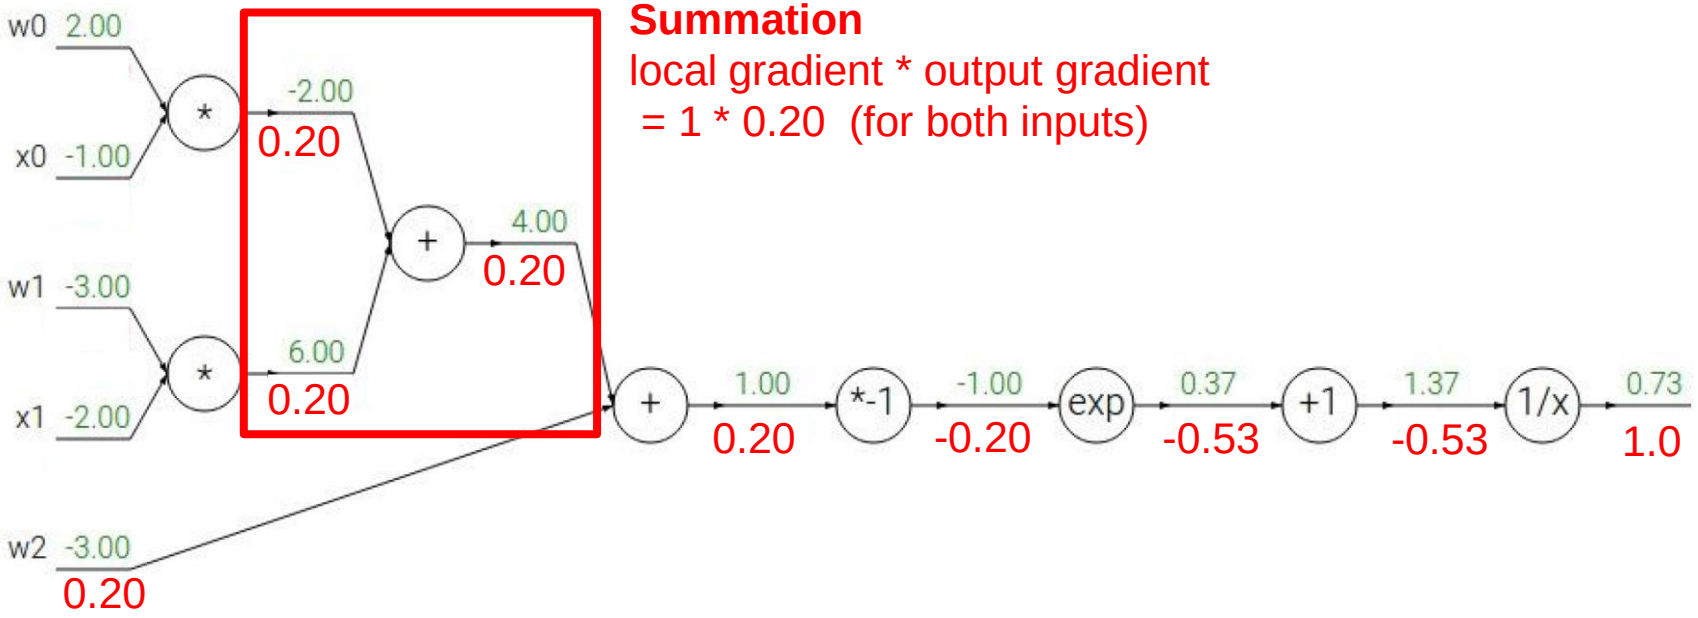

Summation
local gradient * output gradient
= 1 * 0.20 (for both inputs)

$$f(x) = e^x \rightarrow \frac{df}{dx} = e^x \qquad f(x) = \frac{1}{x} \rightarrow \frac{df}{dx} = -1/x^2$$

$$f_a(x) = ax \rightarrow \frac{df}{dx} = a \qquad f_c(x) = c + x \rightarrow \frac{df}{dx} = 1$$

Another example

$$f(w,x) = \frac{1}{1+e^{-(w_0x_0+w_1x_1+w_2)}}$$

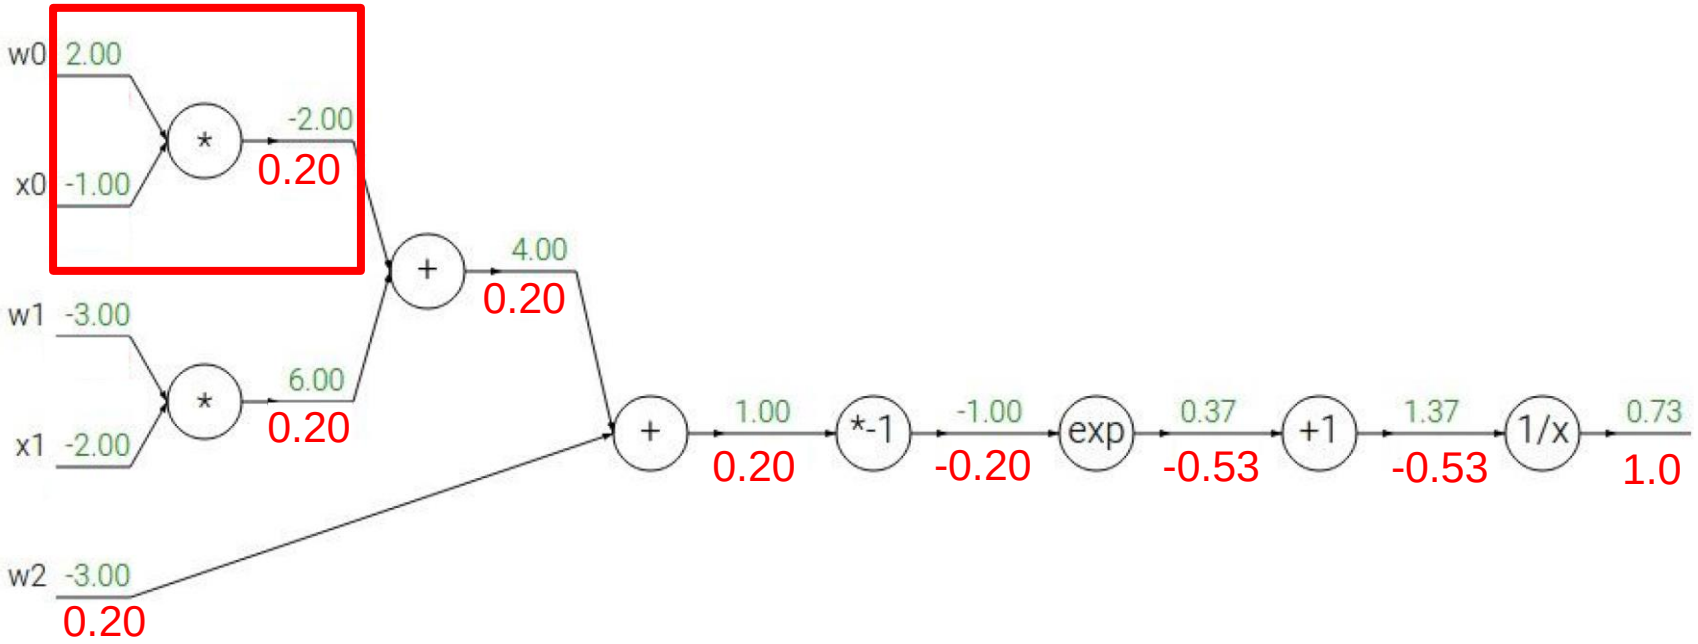

$$f(x) = e^x \rightarrow \frac{df}{dx} = e^x$$

$$f_a(x) = ax \rightarrow \frac{df}{dx} = a$$

$$f(x) = \frac{1}{x} \rightarrow \frac{df}{dx} = -1/x^2$$

$$f_c(x) = c + x \rightarrow \frac{df}{dx} = 1$$

Another example

$$f(w,x) = \frac{1}{1+e^{-(w_0x_0+w_1x_1+w_2)}}$$

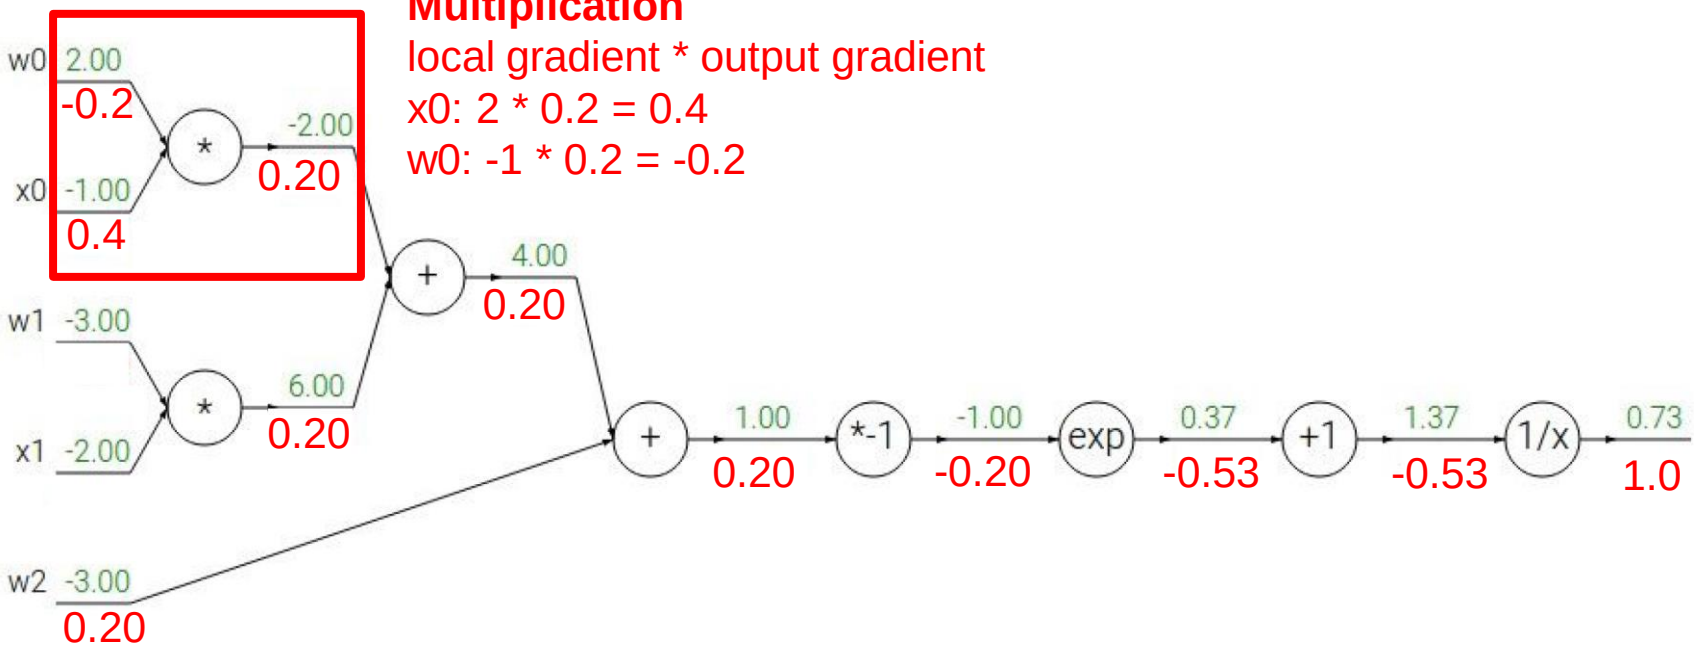

$$f(x) = e^x \rightarrow \frac{df}{dx} = e^x \qquad f(x) = \frac{1}{x} \rightarrow \frac{df}{dx} = -1/x^2$$

$$f_a(x) = ax \rightarrow \frac{df}{dx} = a \qquad f_c(x) = c + x \rightarrow \frac{df}{dx} = 1$$

Another example

$$f(w,x) = \frac{1}{1+e^{-(w_0x_0+w_1x_1+w_2)}}$$

Multiplication
local gradient * output gradient
x1: -3 * 0.2 = -0.6
w1: -2 * 0.2 = -0.4

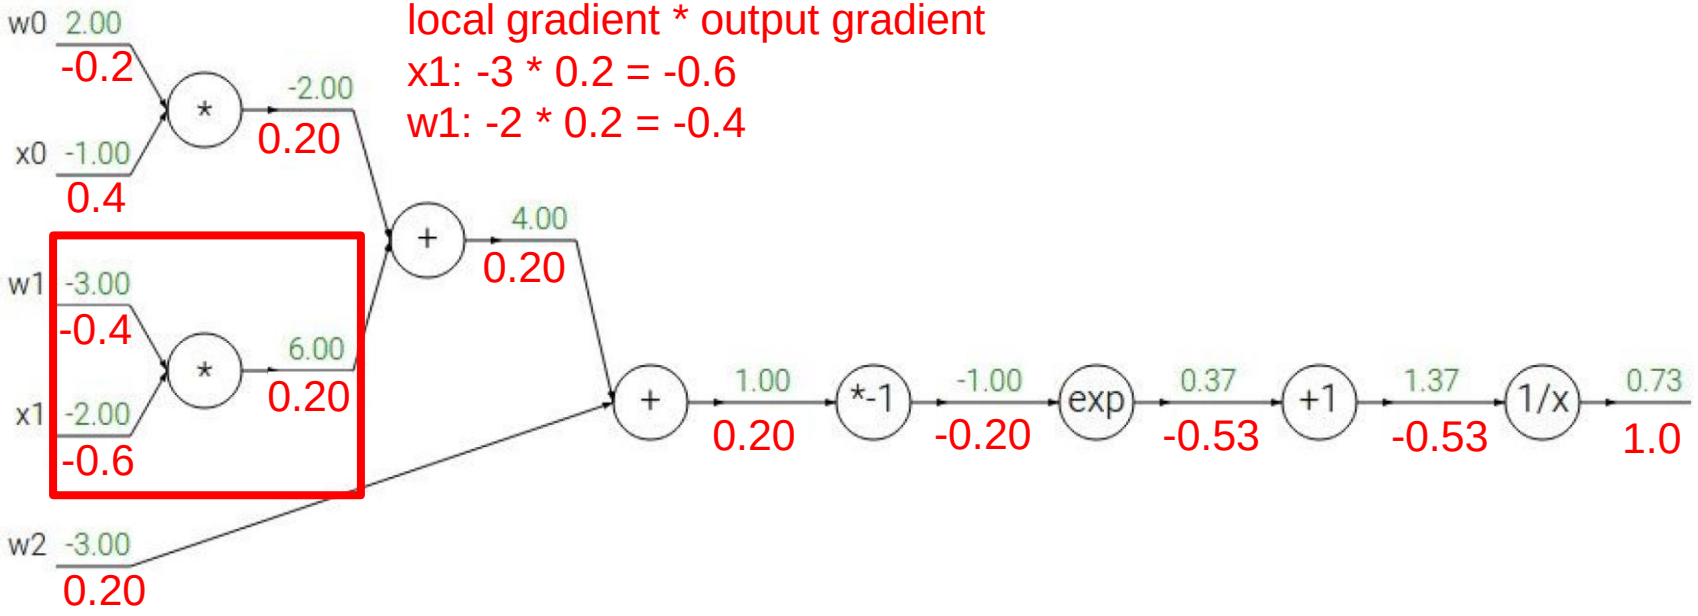

$$f(x) = e^x \rightarrow \frac{df}{dx} = e^x \qquad f(x) = \frac{1}{x} \rightarrow \frac{df}{dx} = -1/x^2$$

$$f_a(x) = ax \rightarrow \frac{df}{dx} = a \qquad f_c(x) = c + x \rightarrow \frac{df}{dx} = 1$$

Another example

$$f(w,x) = \frac{1}{1+e^{-(w_0x_0+w_1x_1+w_2)}}$$

$$\sigma(x) = \frac{1}{1+e^{-x}}$$

sigmoid function

$$\frac{d\sigma(x)}{dx} = \frac{e^{-x}}{(1+e^{-x})^2} = \left(\frac{1+e^{-x}-1}{1+e^{-x}}\right)\left(\frac{1}{1+e^{-x}}\right) = (1-\sigma(x))\,\sigma(x)$$

sigmoid's derivative

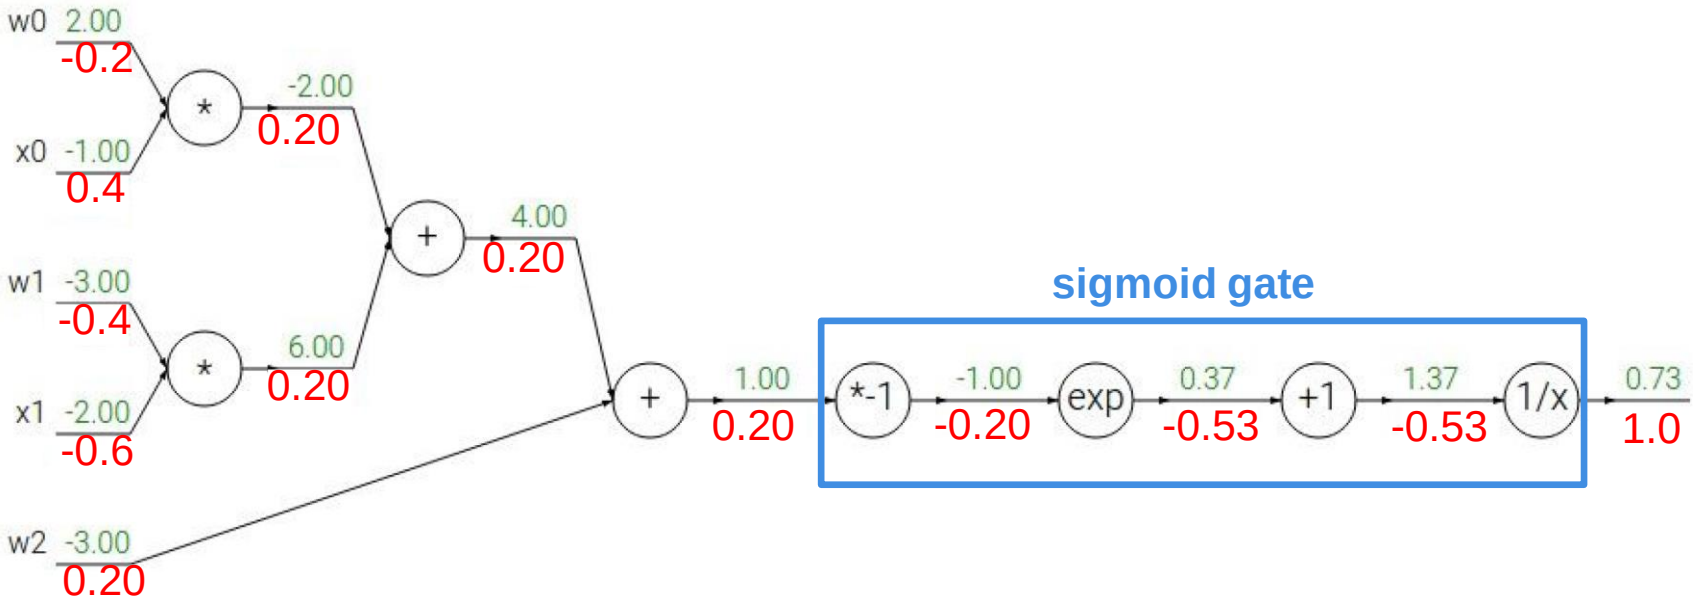

Another example

$$f(w,x) = \frac{1}{1+e^{-(w_0x_0+w_1x_1+w_2)}}$$

$$\sigma(x) = \frac{1}{1+e^{-x}}$$ sigmoid function

$$\frac{d\sigma(x)}{dx} = \frac{e^{-x}}{(1+e^{-x})^2} = \left(\frac{1+e^{-x}-1}{1+e^{-x}}\right)\left(\frac{1}{1+e^{-x}}\right) = (1-\sigma(x))\,\sigma(x)$$ **sigmoid's derivative**

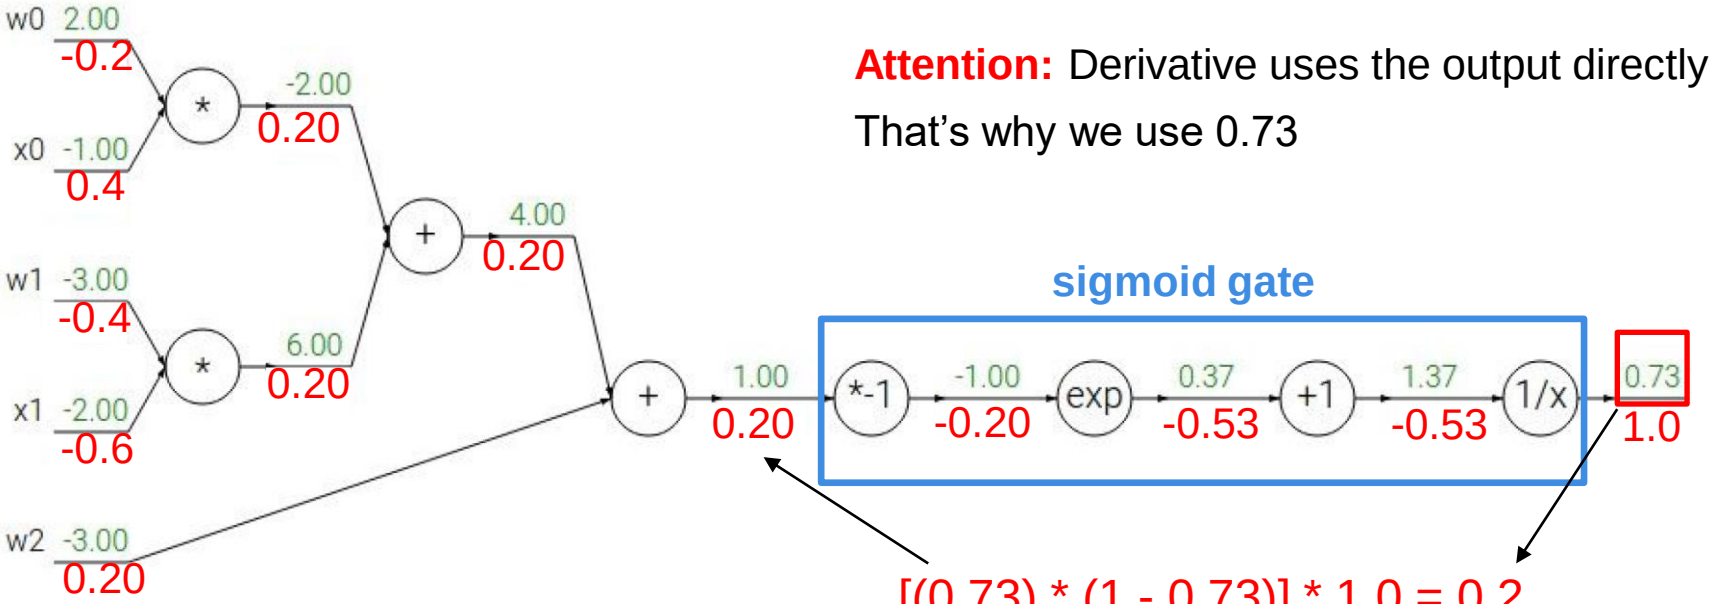

Attention: Derivative uses the output directly

That's why we use 0.73

[(0.73) * (1 - 0.73)] * 1.0 = 0.2

local gradient output gradient

Example of weight update

Gradients, when applied to weights, decrease the loss at each iteration.

Let's say that correct value for a training sample is 0.43 whereas the neuron gives 0.73 at the beginning (loss is 0.3).

Update in the 1st round (we subtract the gradients):

weights: w_0=2.06 w_1=-2.88 w_2=-3.06 function value: 0.66

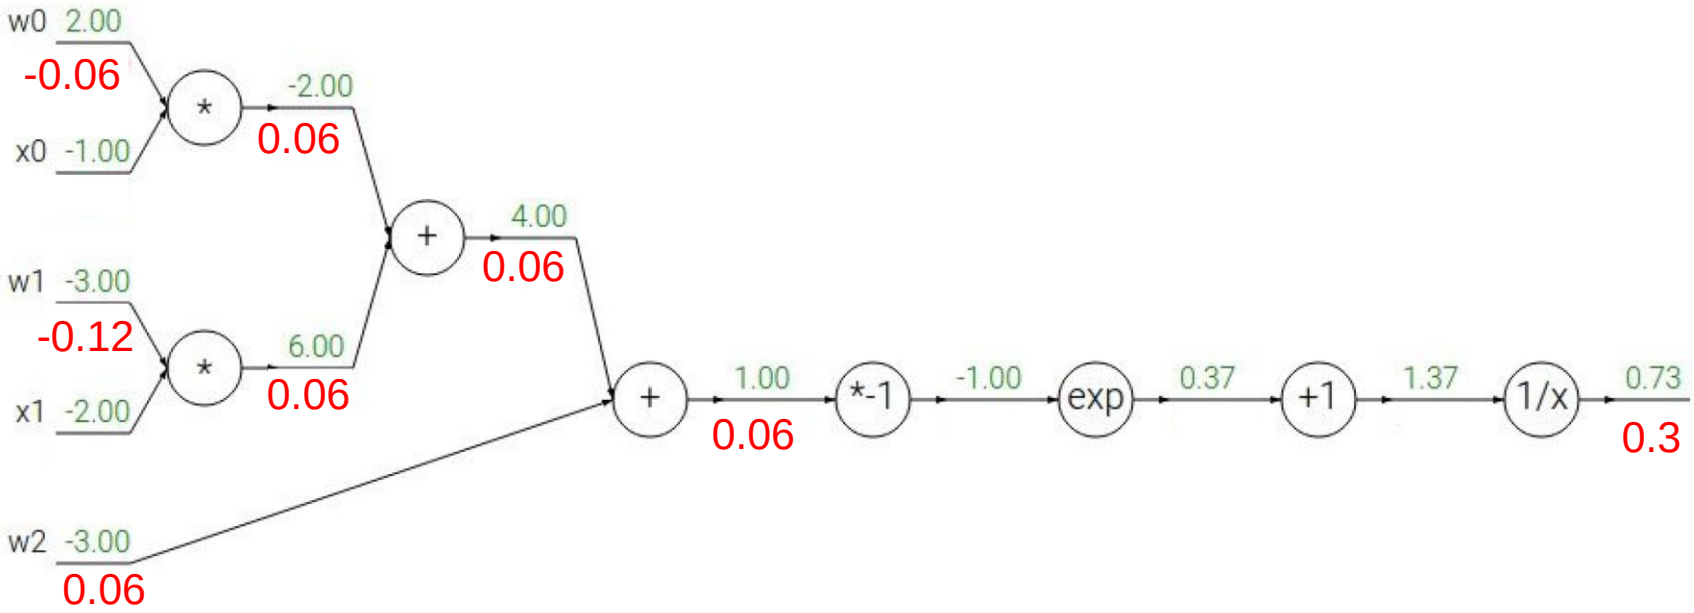

Example of weight update

Gradients, when applied to weights, decreases the loss at each iteration.

Let's say that correct value for a training sample is 0.43 whereas the neuron gives 0.73 at the beginning (loss is 0.3).

Update in the,

1st round, weights: w_0=2.06 w_1=-2.88 w_2=-3.06 function value: 0.66

2nd round, weights: w_0=2.11 w_1=-2.78 w_2=-3.11 function value: 0.58

3rd round, weights: w_0=2.14 w_1=-2.70 w_2=-3.15 function value: 0.53

..

..

..

19th round, weights: w_0=2.21 w_1=-2.57 w_2=-3.21 function value: **0.43**

Forward Propagation Refresher

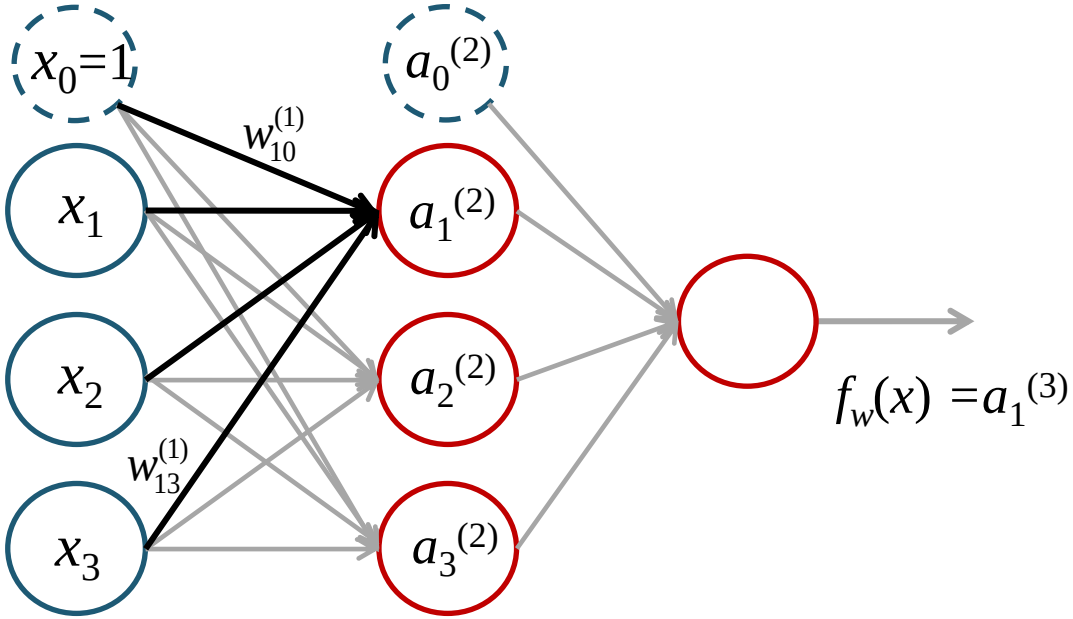

$$z_1^{(2)} = w_{10}^{(1)} + w_{11}^{(1)} x_1 + w_{12}^{(1)} x_2 + w_{13}^{(1)} x_3$$

$$a_1^{(2)} = g(z_1^{(2)})$$

Backpropagation: multiple layers

Need to compute $\frac{\partial E}{\partial w_{ij}^{(l)}}$ for weight update

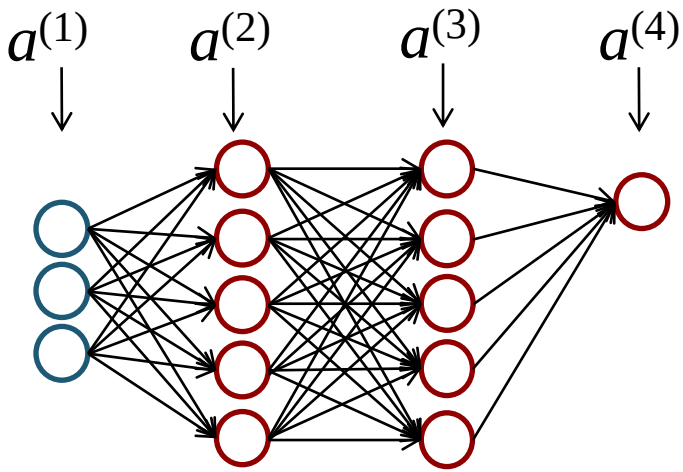

Backpropagation: multiple layers

$\delta_j^{(l)} = \dfrac{\partial E}{\partial z_j^{(l)}}$, aka node delta, is the error of unit j in layer l .

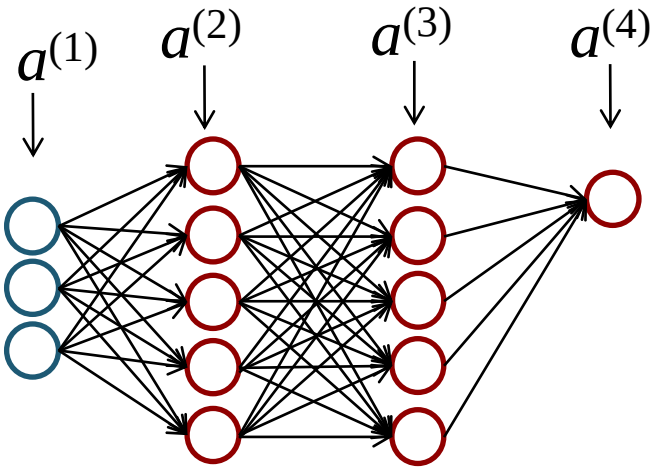

For output unit (layer 4 here) error is:

$$\delta_1^{(4)} = \underbrace{(a_1^{(4)} - y)}_{\substack{\textit{loss,} \\ \textit{output gradient}}} \cdot \underbrace{\frac{\partial a_1^{(4)}}{\partial z_1^{(4)}}}_{\substack{\textit{local} \\ \textit{gradient}}}$$

$$\delta_1^{(4)} = (a_1^{(4)} - y) \cdot \underbrace{a_1^{(4)} \cdot (1 - a_1^{(4)})}_{\textit{local gradient}}$$

Note: We assumed that we used sigmoid in the output unit $\delta_1^{(4)}$

Backpropagation: example

$$\delta_1^{(3)} = \frac{\partial E}{\partial z_1^{(3)}} = \delta_1^{(4)} \cdot w_{11}^{(3)} \cdot a_1^{(3)} \cdot (1 - a_1^{(3)})$$

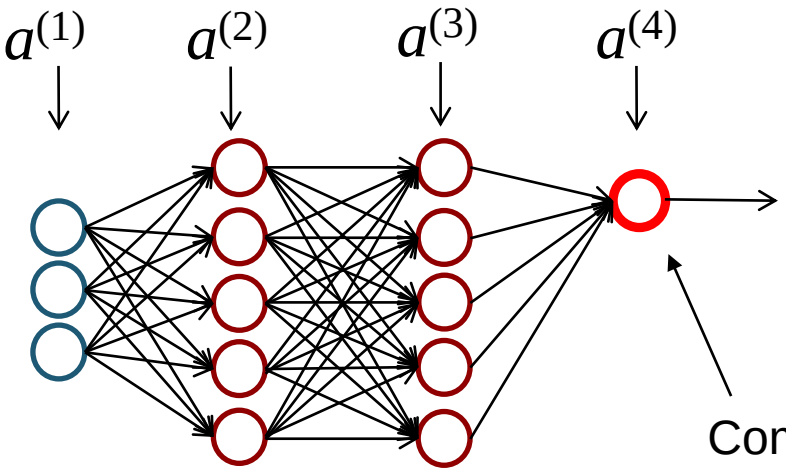

Computed previously

$$\delta_1^{(4)} = (a_1^{(4)} - y) \cdot a_1^{(4)} \cdot (1 - a_1^{(4)})$$

We ‘backpropagate’ the error.

Backpropagation: example

$$\delta_1^{(3)} = \frac{\partial E}{\partial z_1^{(3)}} = \delta_1^{(4)} \cdot \boxed{w_{11}^{(3)}} \cdot a_1^{(3)} \cdot (1 - a_1^{(3)})$$

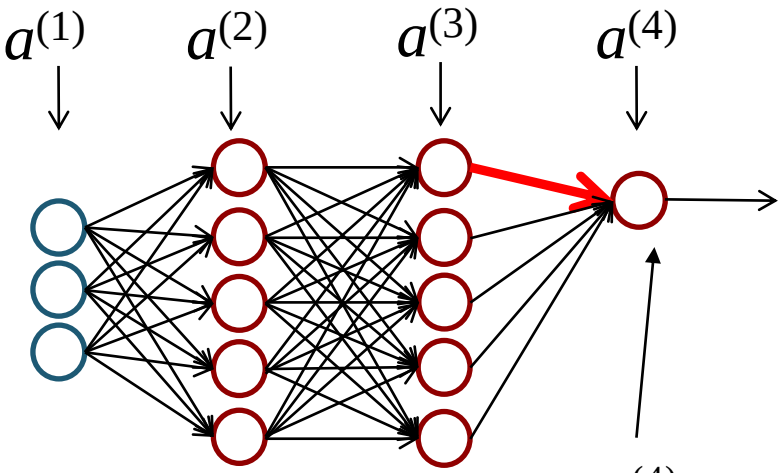

$$\frac{\partial z_1^{(4)}}{\partial a_1^{(3)}} = \frac{\partial (w_{11}^{(3)} a_1^{(3)} + w_{12}^{(3)} a_2^{(3)} + w_{13}^{(3)} a_3^{(3)} + ..)}{\partial a_1^{(3)}}$$

Backpropagation: example

$$\delta_1^{(3)} = \frac{\partial E}{\partial z_1^{(3)}} = \delta_1^{(4)} \cdot w_{11}^{(3)} \cdot a_1^{(3)} \cdot (1 - a_1^{(3)})$$

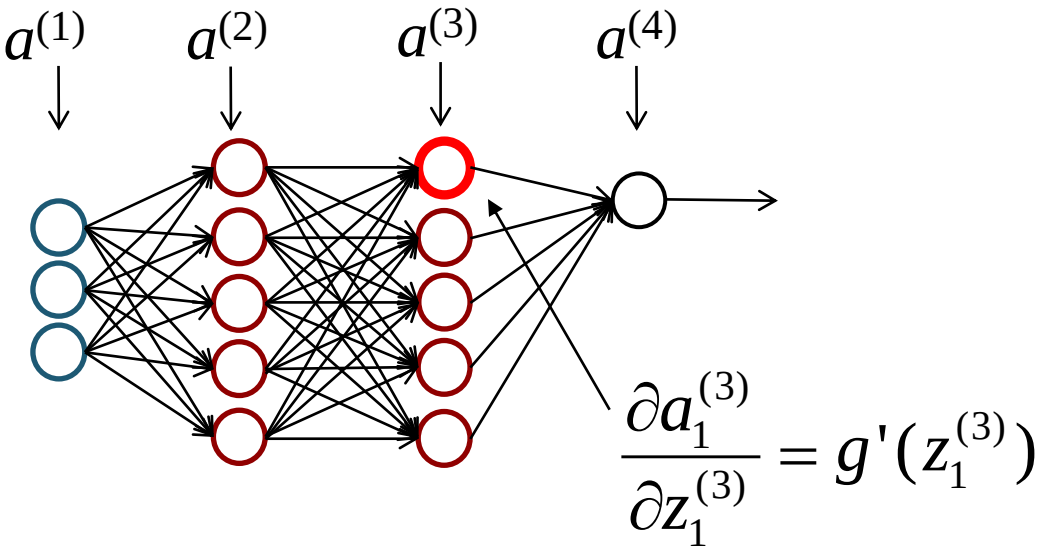

Backpropagation: weight gradient

$$\frac{\partial(w_{11}^{(3)}a_1^{(3)} + w_{12}^{(3)}a_2^{(3)} + w_{13}^{(3)}a_3^{(3)} + ..)}{\partial w_{11}^{(3)}}$$

E.g. $$\frac{\partial E}{\partial w_{11}^{(3)}} = \frac{\partial E}{\partial z_1^{(4)}} \frac{\partial z_1^{(4)}}{\partial w_{11}^{(3)}} = \delta_1^{(4)} a_1^{(3)} = (a_1^{(4)} - y) \cdot a_1^{(4)} \cdot (1 - a_1^{(4)}) \cdot a_1^{(3)}$$

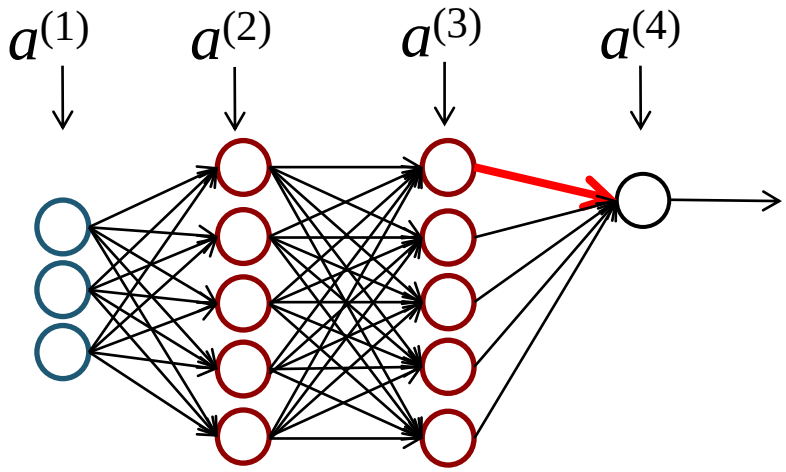

Backpropagation: weight gradient

$$\frac{\partial(w_{11}^{(2)}a_1^{(2)} + w_{12}^{(2)}a_2^{(2)} + w_{13}^{(2)}a_3^{(2)} + ..)}{\partial w_{13}^{(2)}}$$

Another E.g.

$$\frac{\partial E}{\partial w_{13}^{(2)}} = \frac{\partial E}{\partial z_1^{(3)}} \frac{\partial z_1^{(3)}}{\partial w_{13}^{(2)}} = \delta_1^{(3)} a_3^{(2)} = (a_1^{(4)} - y) \cdot a_1^{(4)} \cdot (1 - a_1^{(4)}) \cdot w_{11}^{(3)} \cdot a_1^{(3)} \cdot (1 - a_1^{(3)}) \cdot a_3^{(2)}$$

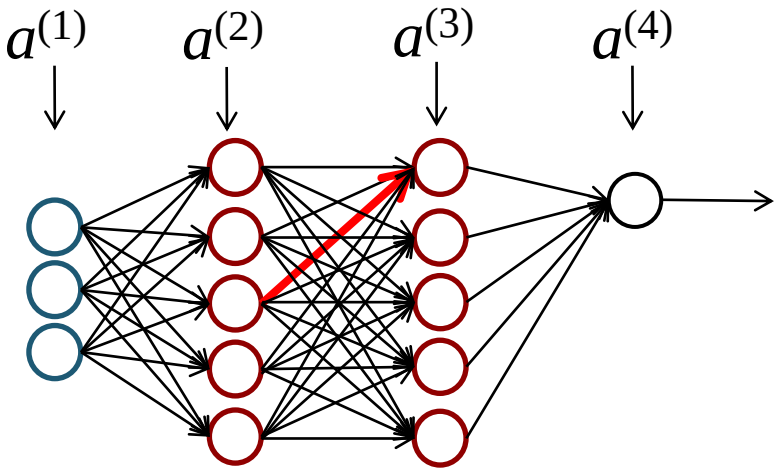

An example of backpropagation with actual numbers

Please see the example given in https://mattmazur.com/2015/03/17/a-step-by-step-backpropagation-example/comment-page-5/#comments

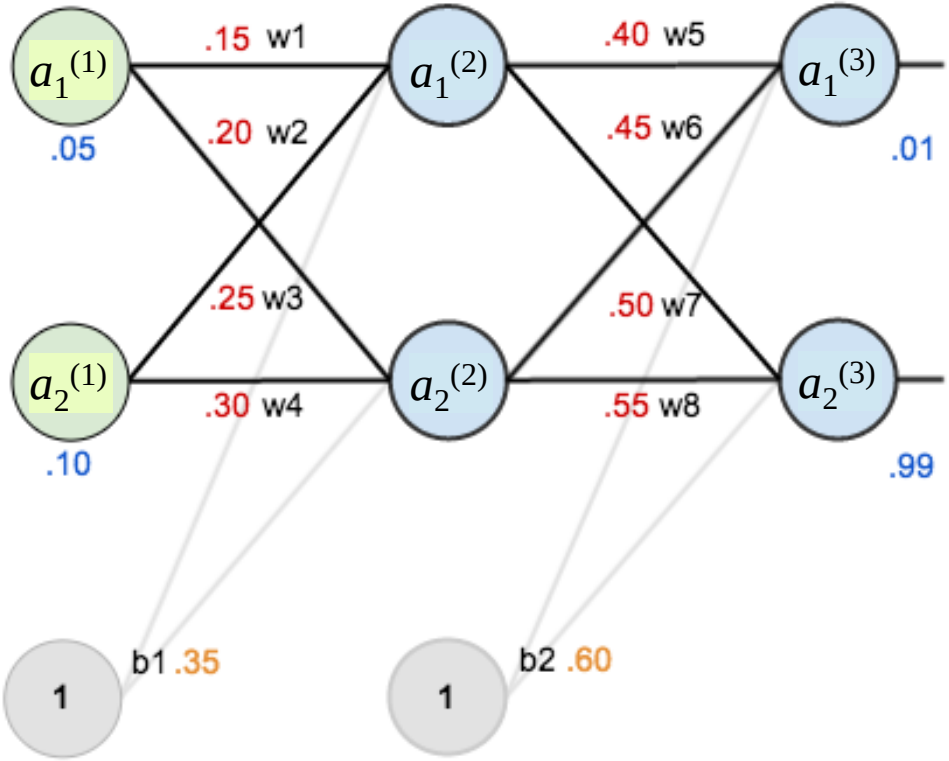

Forward pass

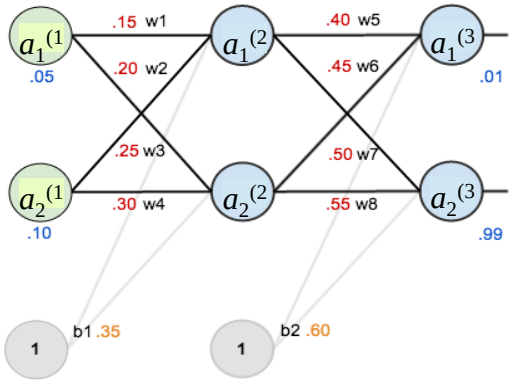

Backward pass (1)

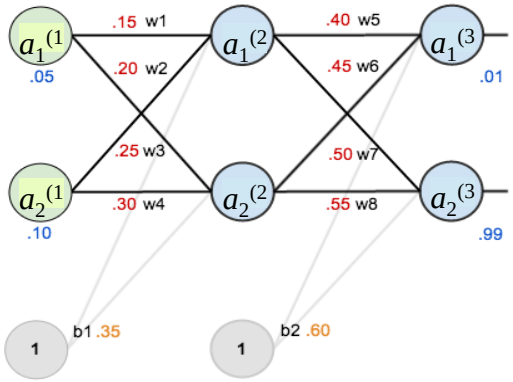

Backward pass (2)

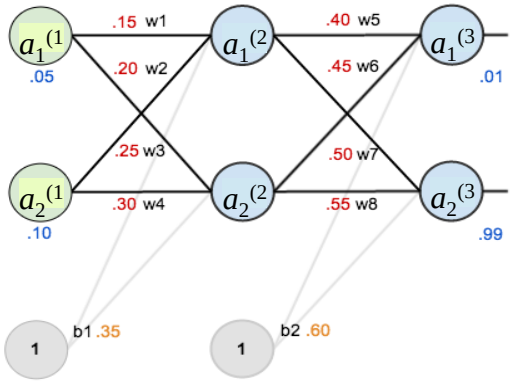

Putting it together

Training a neural network

1. Randomly initialize weights
2. Implement forward propagation to get $a^{(L)}$ where L: last layer
3. Implement code to compute loss, i.e. cost function E
4. Implement backpropagation to compute partial derivatives: $\frac{\partial E}{\partial w_{ij}^{(l)}}$
5. Use gradient descent or an advanced optimization method with backpropagation and try to minimize E as a function of parameters

Note: *Batch size* means how many samples you see to do one cycle of weight update. If you compute E over 8 samples (images) and you update all the weights based on this E, then your *batch size* = 8.

Next

Next week, we will see the whole procedure to train neural networks.