

# CENG 506 Deep Learning

## Lecture 9 – Explainability in Deep Learning

# Why Explainability?


- Deep learning is usually more successful than other ML methods like logistic regression, k-NN, decision trees etc. but its explainability level is low.
- It becomes more important to be able to explain ML to various stakeholders (customers, managers etc)
- Regulators/authorities ask more information on how ML models make their decisions in various domains (healthcare, insurance, finance etc)

# What is going on inside ConvNets?

This image is CC0 public domain



Input Image:  
3 x 224 x 224




What are the intermediate features looking for?


# First Layers

We can easily visualize them (as  $F \times F \times 3$  RGB images)


They are sensitive to oriented edges, dots, dominant or opposing colors etc.




AlexNet:  
 $64 \times 3 \times 11 \times 11$



ResNet-18:  
 $64 \times 3 \times 7 \times 7$



ResNet-101:  
 $64 \times 3 \times 7 \times 7$



DenseNet-121:  
 $64 \times 3 \times 7 \times 7$

# Visualizing Intermediate Layers

Weights:



Further layers are less interpretable

layer 2 weights

$20 \times 16 \times 7 \times 7$

Weights:





layer 3 weights

$20 \times 20 \times 7 \times 7$

Source: [ConvNetJS](#)  
[CIFAR-10 demo](#)



# Maximally Activating Patches




Pick a layer and a channel;  
e.g. conv5 is 128 x 13 x 13,  
pick 17th channel

Run many images through the  
network, record values of  
chosen channel

Visualize image patches that  
correspond to maximal  
activations

# Intermediate features via backprop





Pick a single intermediate neuron,  
e.g. one value in 128 x 13 x 13 conv5 feature map.

Compute gradient of neuron value with respect to image pixels (only backprop positive gradients - **guided backprop**).






# Guided backprop



Zeiler and Fergus, "Visualizing and Understanding Convolutional Networks", ECCV 2014  
Springenberg et al, "Striving for Simplicity: The All Convolutional Net", ICLR Workshop 2015



# Backprop at different level features



# Fooling Images / Adversarial Examples

- (1) Start from an arbitrary image
- (2) Pick an arbitrary class
- (3) Modify the image to maximize the class
- (4) Repeat until network is fooled

# Fooling Images / Adversarial Examples

African elephant




koala



Difference



10x Difference



schooner




iPod



Difference

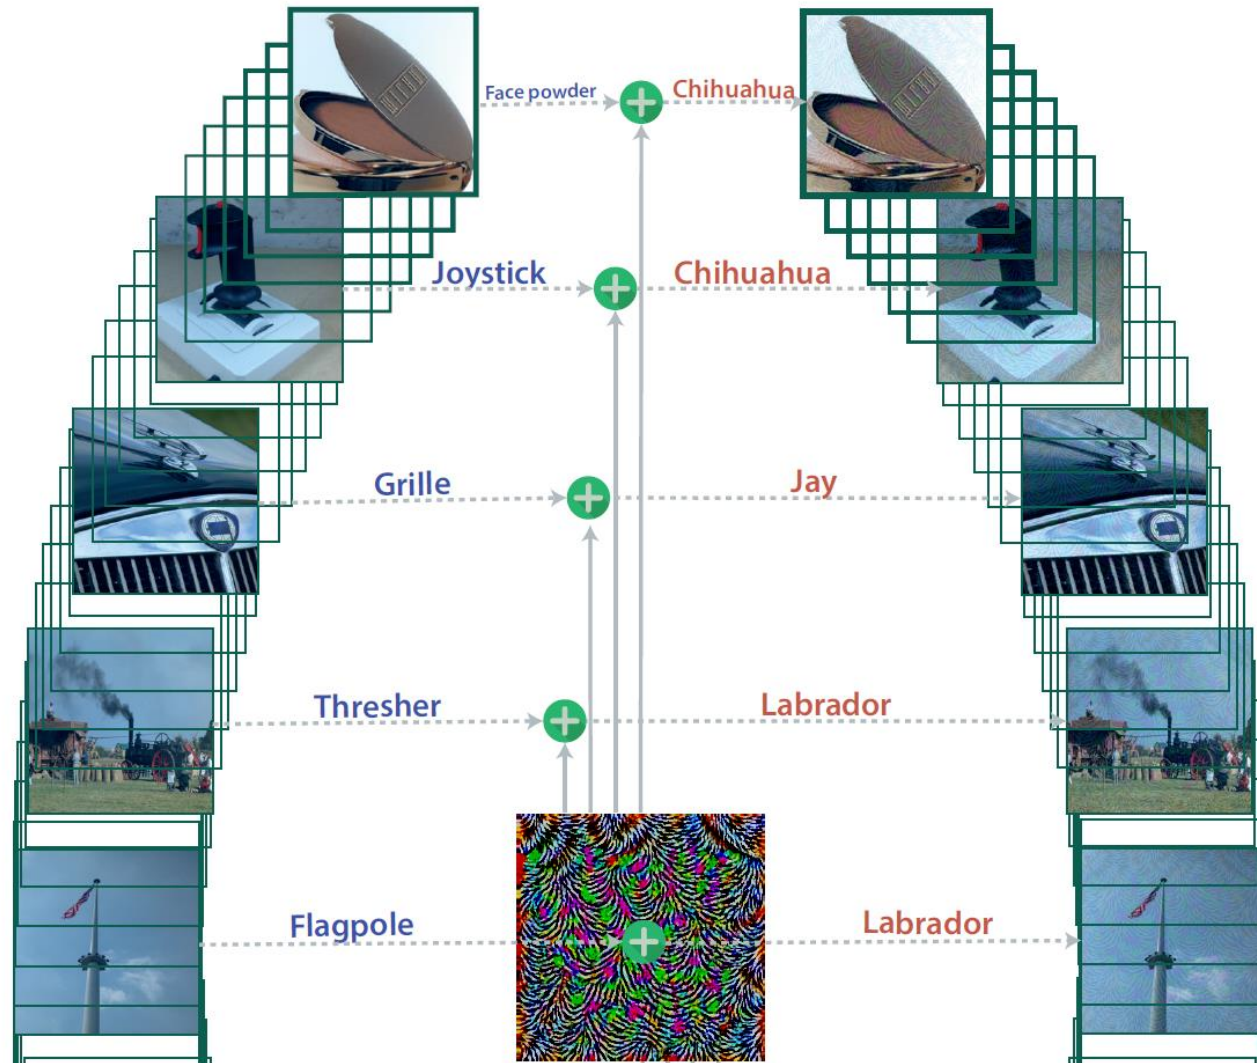


10x Difference






# Training Perturbations to Fool Networks






# Training Perturbations to Fool Networks




(a) CaffeNet



(b) VGG-F




(c) VGG-16



(d) VGG-19



(e) GoogLeNet



(f) ResNet-152

# Training Perturbations to Fool Networks

	VGG-F	CaffeNet	GoogLeNet	VGG-16	VGG-19	ResNet-152
VGG-F	<b>93.7%</b>	71.8%	48.4%	42.1%	42.1%	47.4 %
CaffeNet	74.0%	<b>93.3%</b>	47.7%	39.9%	39.9%	48.0%
GoogLeNet	46.2%	43.8%	<b>78.9%</b>	39.2%	39.8%	45.5%
VGG-16	63.4%	55.8%	56.5%	<b>78.3%</b>	73.1%	63.4%
VGG-19	64.0%	57.2%	53.6%	73.5%	<b>77.8%</b>	58.0%
ResNet-152	46.3%	46.3%	50.5%	47.0%	45.5%	<b>84.0%</b>

Generalizability of the universal perturbations across different networks. **The percentages indicate the fooling rates (for ImageNet data).** The rows indicate the architecture for which the universal perturbations is computed, and the columns indicate the architecture for which the fooling rate is reported. Not surprisingly, networks are best fooled if perturbation is computed on itself.

# Why Test Accuracy is not Enough to trust on ConvNets?

Test accuracy may not capture critical issues

- Bad data
- Biases
- Poor performance in critical cases



# Why Test Accuracy is not Enough?

E.g. Train a neural network to predict wolf vs. husky



Husky



Wolf

Some results:



Predicted: **wolf**  
True: **wolf**




Predicted: **husky**  
True: **husky**



Predicted: **wolf**  
True: **wolf**



Predicted: **wolf**  
True: **husky**



Predicted: **husky**  
True: **husky**



Predicted: **wolf**  
True: **wolf**

Source:

Why should I trust you?: Explaining the Predictions of Any Classifier. Ribeiro, Singh & Guestrin KDD'16




# Why Test Accuracy is not Enough?

## Explanations for neural network prediction



Predicted: **wolf**  
True: **wolf**



Predicted: **husky**  
True: **husky**



Predicted: **wolf**  
True: **wolf**



Predicted: **wolf**  
True: **husky**



Predicted: **husky**  
True: **husky**



Predicted: **wolf**  
True: **wolf**

Source:

Why should I trust you?: Explaining the Predictions of Any Classifier. Ribeiro, Singh & Guestrin KDD'16

# Why Test Accuracy is not Enough?

## Explanations for neural network prediction




(a) Original Image



(b) Explaining *Electric guitar*



(c) Explaining *Acoustic guitar*



(d) Explaining *Labrador*


Figure 4: Explaining an image classification prediction made by Google's Inception neural network. The top 3 classes predicted are "Electric Guitar" ( $p = 0.32$ ), "Acoustic guitar" ( $p = 0.24$ ) and "Labrador" ( $p = 0.21$ )

Source:


Why should I trust you?: Explaining the Predictions of Any Classifier. Ribeiro, Singh & Guestrin KDD'16

# Gradient Class Activation Mapping (Grad-CAM\*)

Grad-CAM is a technique to visualize where a CNN is looking. It is class-specific, meaning it can produce a separate visualization for every class.



Original Image



Grad-CAM 'Cat'




Grad-CAM 'Dog'

\* Selvaraju et al. "Grad-CAM: Visual Explanations from Deep Networks via Gradient-based Localization." International Journal of Computer Vision 2019.

# Class Activation Mapping (CAM\*)

CAM requires applying global average pooling (GAP) to the final convolutional feature maps, followed by a single fully connected layer that produces the predictions:




\* <https://towardsdatascience.com/grad-cam-2f0c6f3807fe>



# Class Activation Mapping (CAM\*)

For class “cat”, the prediction depends on  $k$  weights ( $w_1, w_2, \dots, w_k$ ). To make a CAM heatmap for “cat”, we perform a weighted sum of the feature maps, using the “cat” weights of the final FC layer:

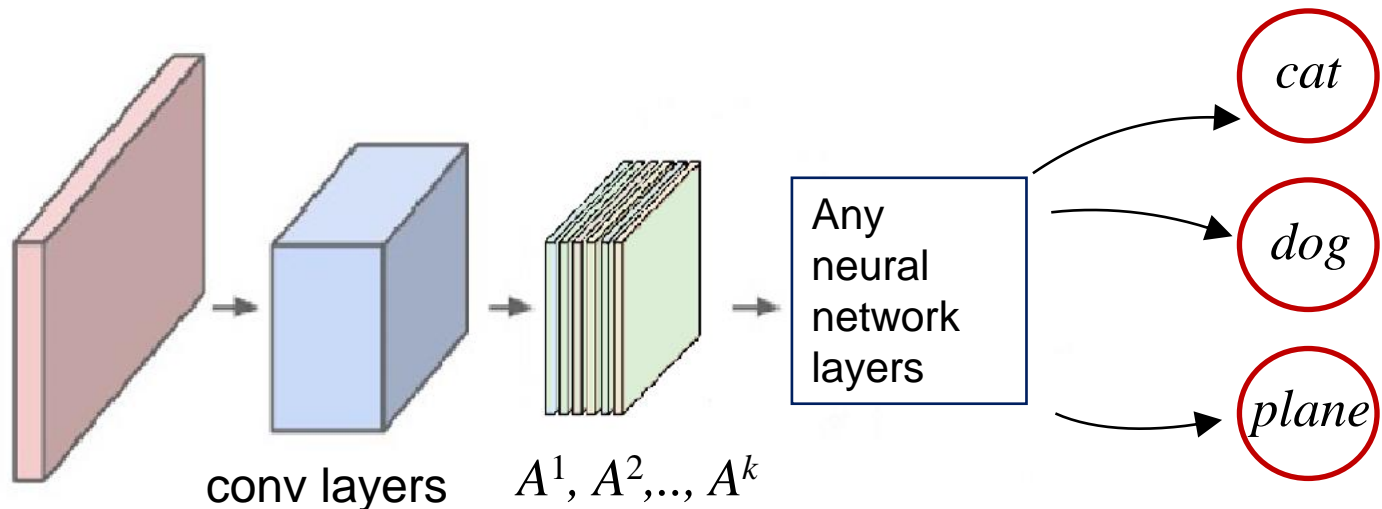


$$CAM^{cat} = w_1 A^1 + w_2 A^2 + \dots + w_k A^k = \sum_i w_i^{cat} A^i$$

\* <https://towardsdatascience.com/grad-cam-2f0c6f3807fe>

# Grad-CAM

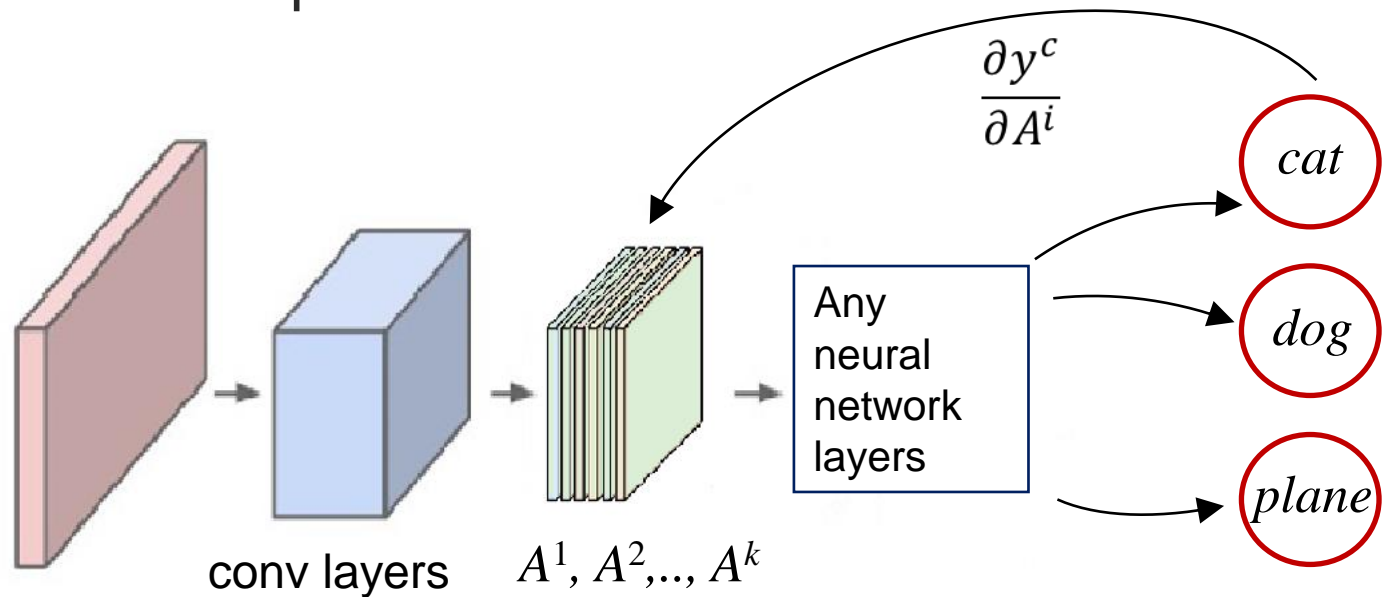
Grad-CAM can use any network between the last conv. layer and the classification layer as long as it is differentiable. We can calculate gradients through any kind of network. I.e. it does not require GAP.



# Grad-CAM

Grad-CAM has three steps:

1) Compute the gradient of scores ( $y^c$ ) w.r.t the last conv layer activation maps



# Grad-CAM

2) Calculate alphas by averaging gradients

Global average pool the gradients over the width x height

$$\alpha_k^c = \overbrace{\frac{1}{Z} \sum_i \sum_j}^{\text{global average pooling}} \underbrace{\frac{\partial y^c}{\partial A_{ij}^k}}_{\text{gradients via backprop}}$$

3) Perform a weighted combination of the activation maps where weights are calculated alphas (an additional RELU)


$$L_{\text{Grad-CAM}}^c = \text{ReLU} \left( \underbrace{\sum_k \alpha_k^c A^k}_{\text{linear combination}} \right)$$



# Grad-CAM

Q. Won't the Grad-CAM heatmap be too small?  
I.e. height and width of the last conv layer.

A. Smoothed while upsampling



# Many variations of CAM\*

Method	What it does
GradCAM	Weight the 2D activations by the average gradient
HiResCAM	Like GradCAM but element-wise multiply the activations with the gradients; provably guaranteed faithfulness for certain models
GradCAM++	Like GradCAM but uses second order gradients
XGradCAM	Like GradCAM but scale the gradients by the normalized activations
AblationCAM	Zero out activations and measure how the output drops (this repository includes a fast batched implementation)
ScoreCAM	Permute the image by the scaled activations and measure how the output drops
EigenCAM	Takes the first principle component of the 2D Activations (no class discrimination, but seems to give great results)
LayerCAM	Spatially weight the activations by positive gradients. Works better especially in lower layers

\* <https://jacobgil.github.io/pytorch-gradcam-book/introduction.html>